

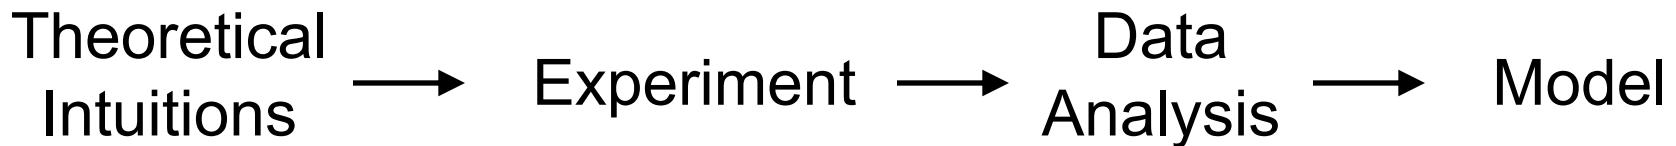
Probabilistic models of cognition

Charles Kemp
CHD Summer school, 2018

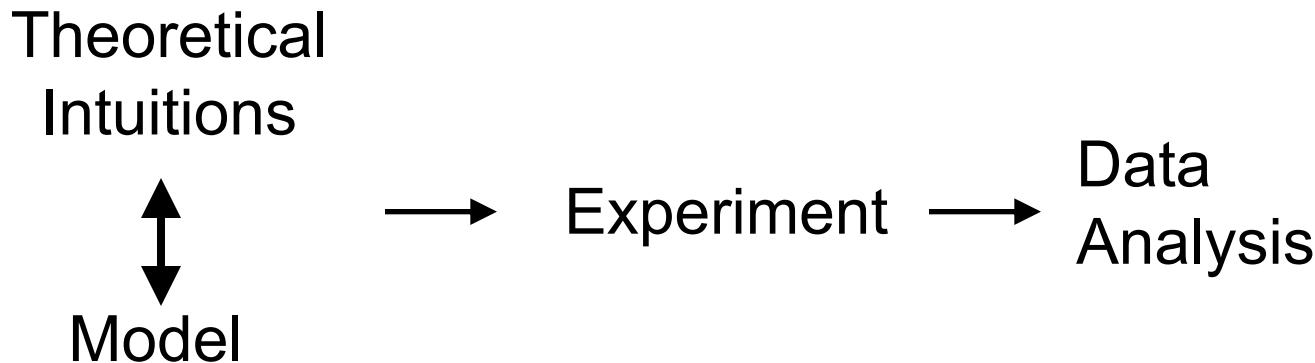
Why build models?

Why build models?

- To help explain some empirical result



- To help formulate a theory



Modeling approaches

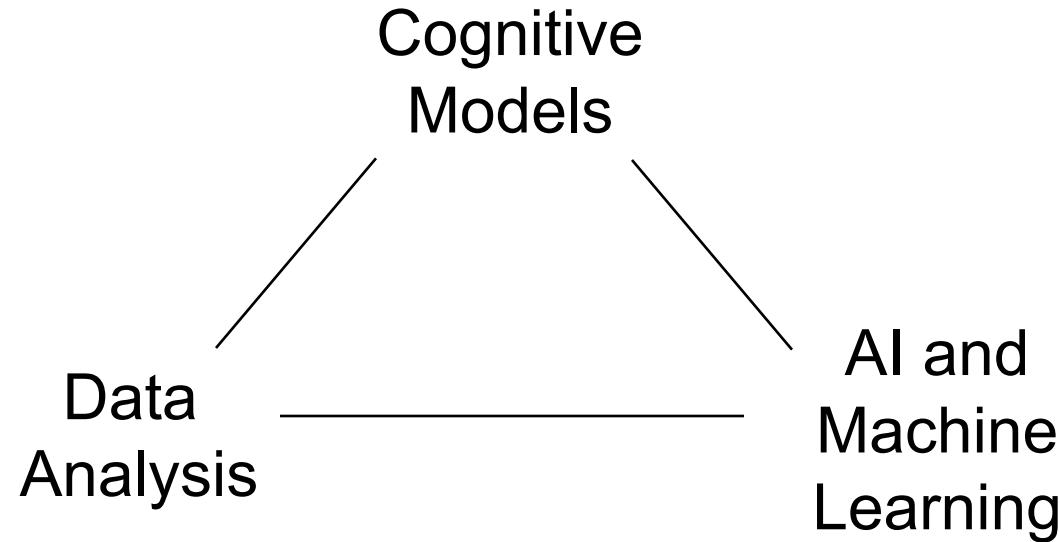
- Formal logic
- Production systems (e.g SOAR, ACT-R)
- Neural networks
- Dynamic systems
- Probabilistic models

Why build *probabilistic* models?

- For us: probabilities are degrees of belief
- Probability theory captures the right way to update degrees of belief

“The theory of probabilities is at bottom only common sense reduced to calculus; it makes us appreciate with exactitude that which exact minds feel by a sort of instinct without being able oft times to give a reason for it” (Laplace)

Applications of probability theory



Probabilistic machine learning and artificial intelligence

Module 1

INTRODUCTION TO BAYESIAN INFERENCE

Inductive inference



Coughing friend

Inductive inference



John has a cold

Coughing friend

Inductive inference



John has emphysema

Coughing friend

Coughing friend

- d : John is coughing
- h_1 : John has a cold
- h_2 : John has emphysema
- h_3 : John has a stomach upset

Posterior
probability

Evidence
(Likelihood)

Prior
knowledge

$$P(h|d) = \frac{P(d|h) P(h)}{P(d)}$$

Coughing friend

- d : John is coughing
- h_1 : John has a cold
- h_2 : John has emphysema
- h_3 : John has a stomach upset

Posterior
probability

Evidence
(Likelihood)

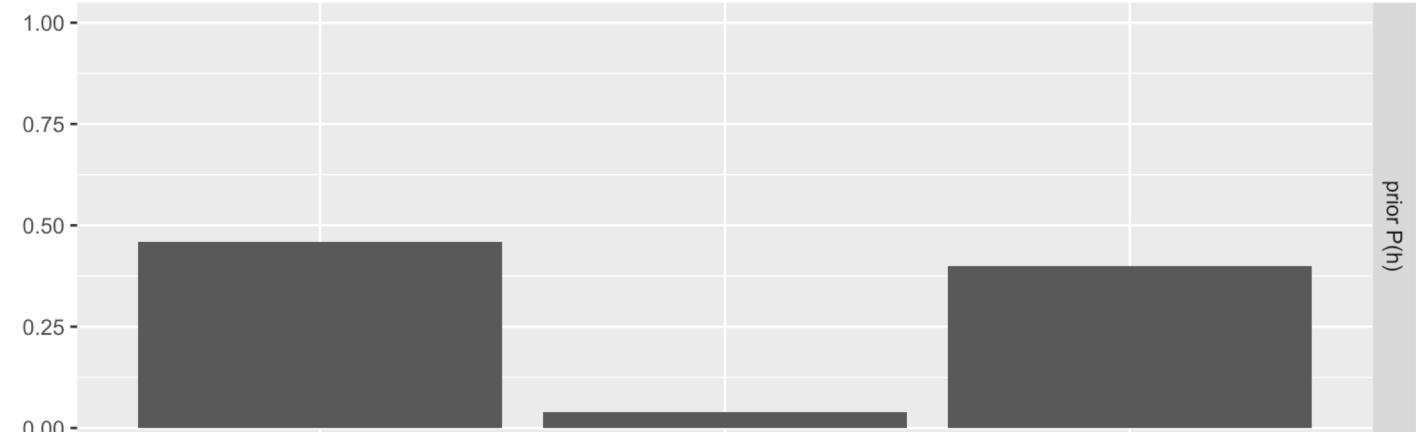
Prior
knowledge

$$P(h|d) \propto P(d|h) P(h)$$

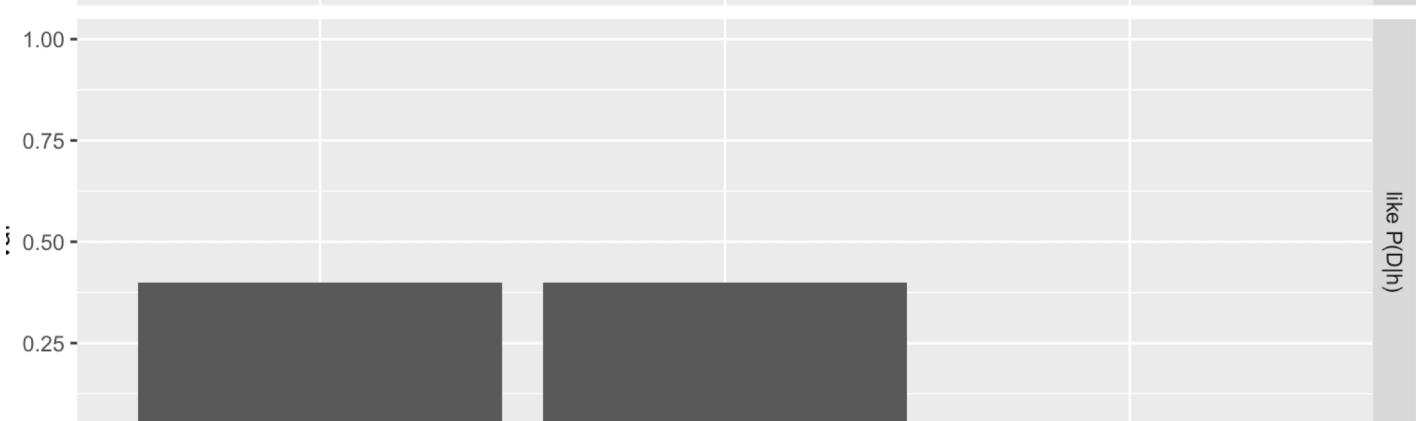
Specifying prior and likelihood

```
h <- c('cold', 'emphysema', 'stomach upset')
p_h <- c(0.46, 0.04, 0.4)
p_d_given_h <- c(0.4, 0.4, 0.05)
```

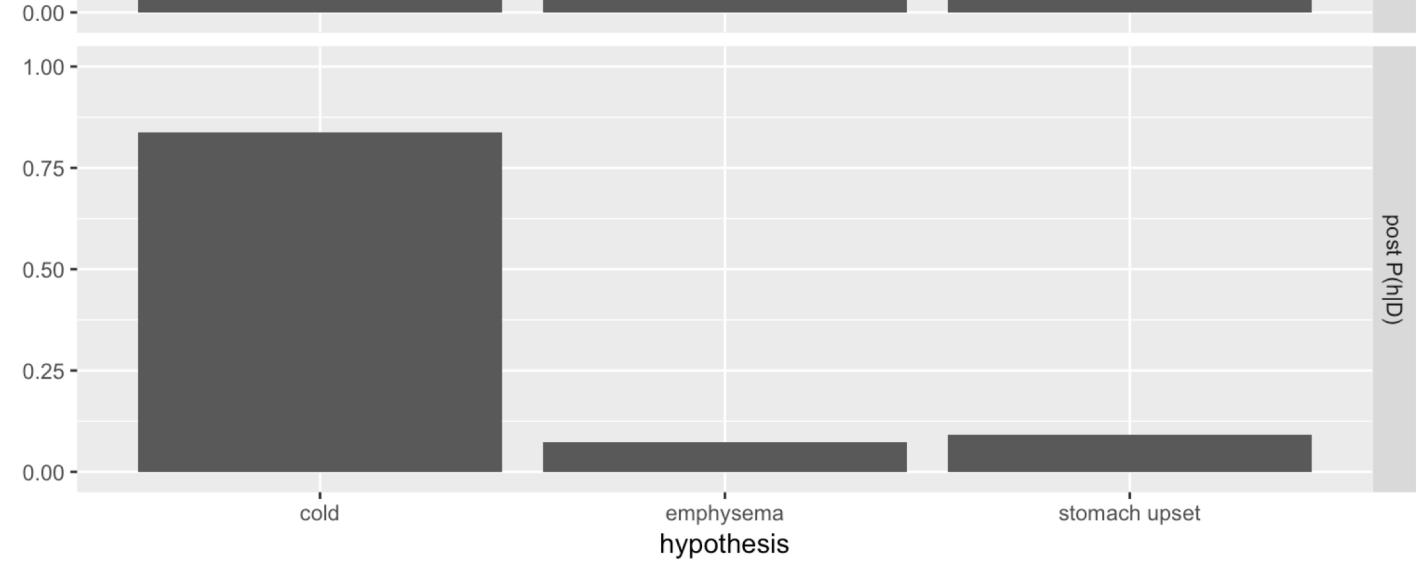
prior



likelihood



posterior



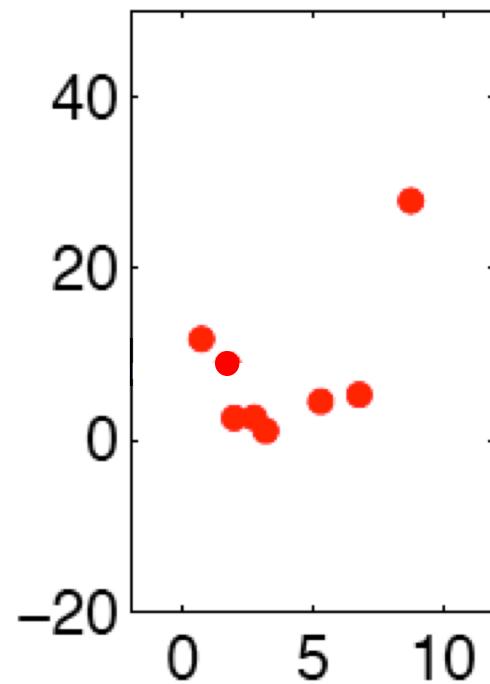
Exercise: Coughing friend

```
h <- c('cold', 'emphysema', 'stomach upset')
p_h <- c(0.46, 0.04, 0.4)
p_d_given_h <- c(0.4, 0.4, 0.05)
```

Inductive inference



Coughing friend



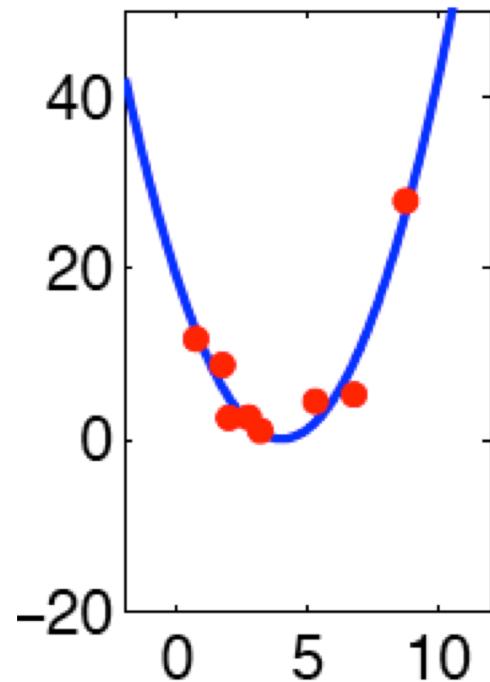
Curve fitting

Inductive inference



John has a cold

Coughing friend



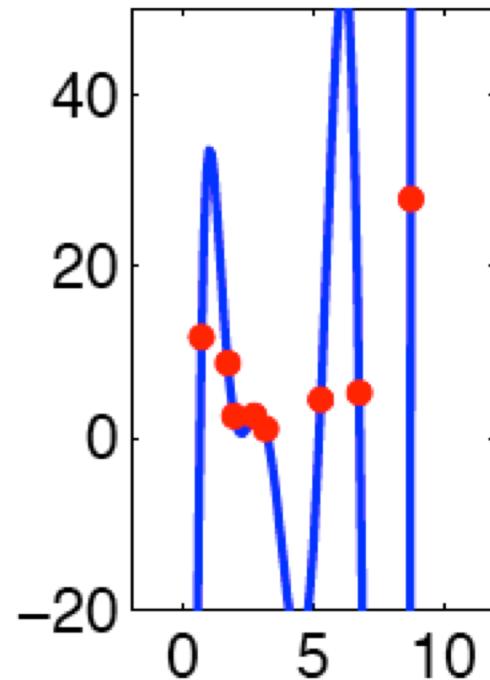
Curve fitting

Inductive inference



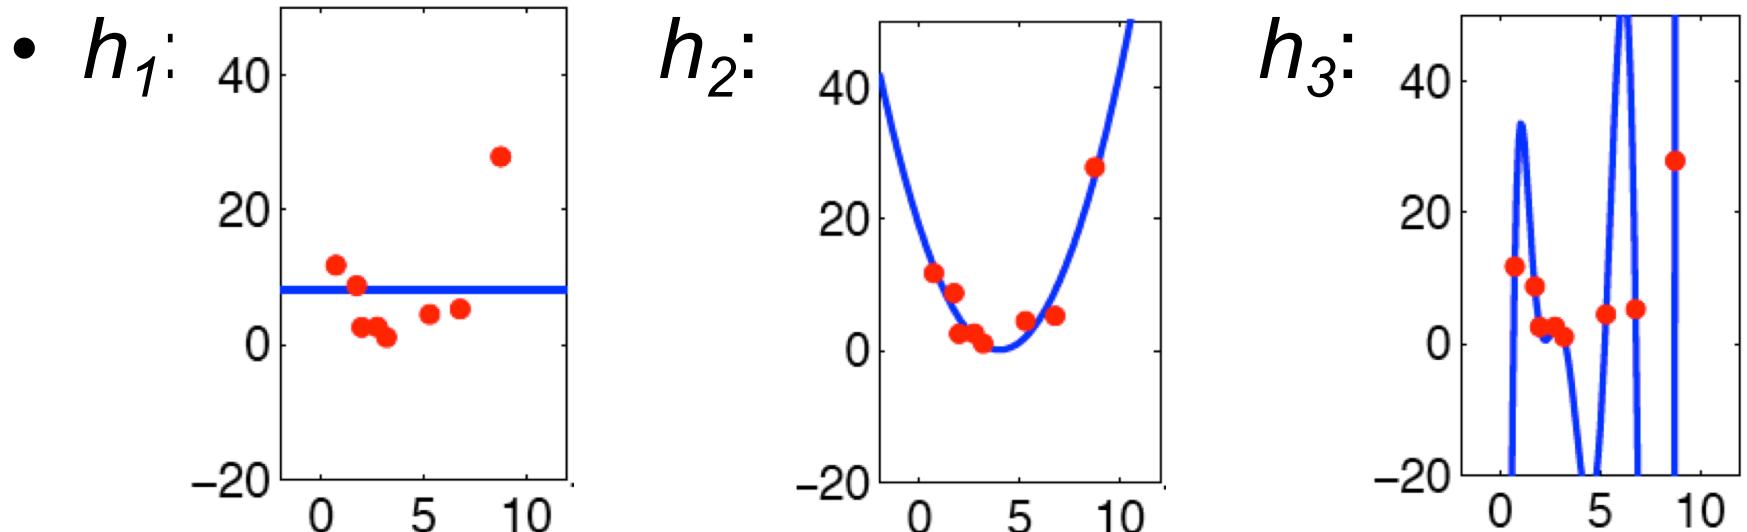
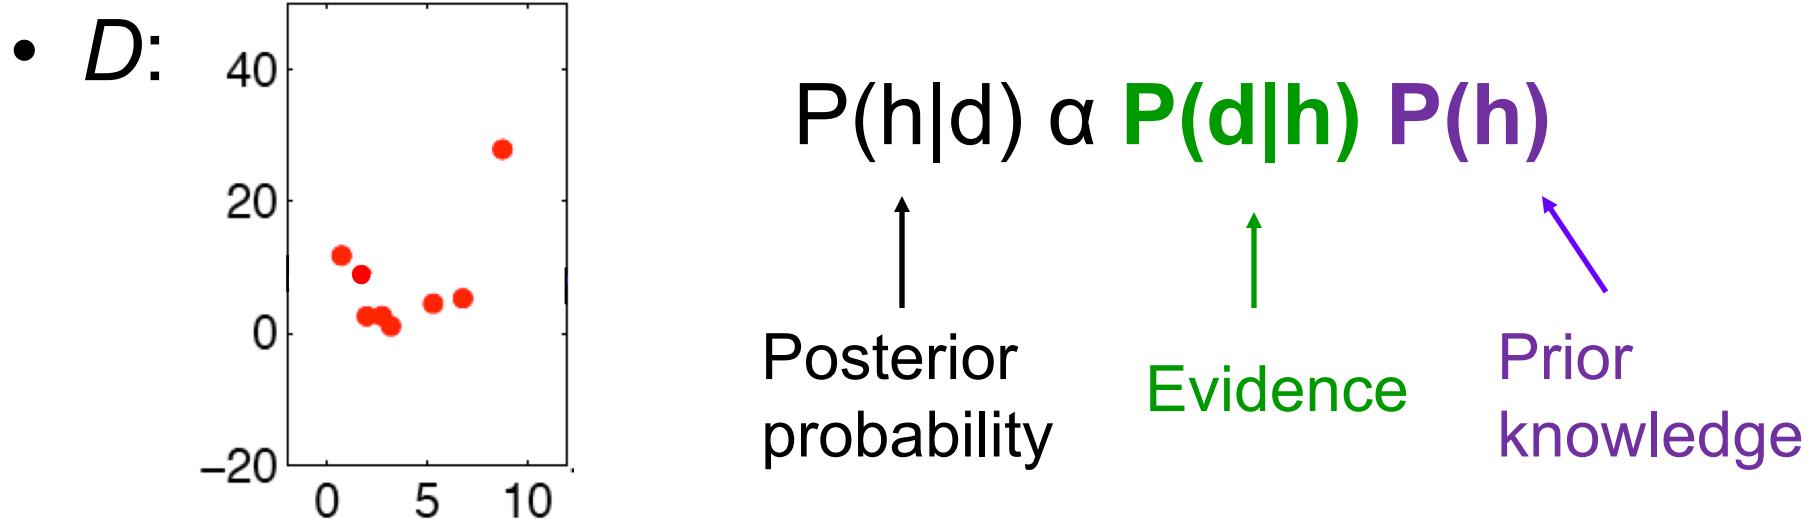
John has emphysema

Coughing friend

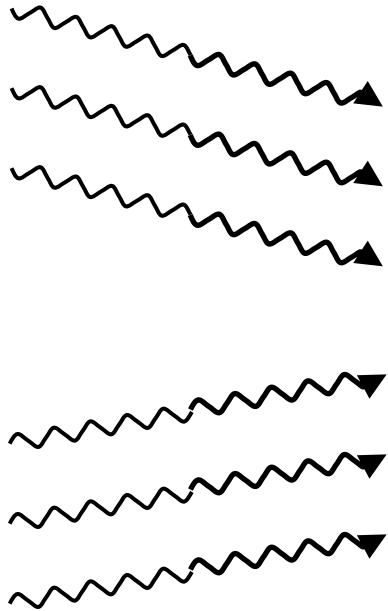


Curve fitting

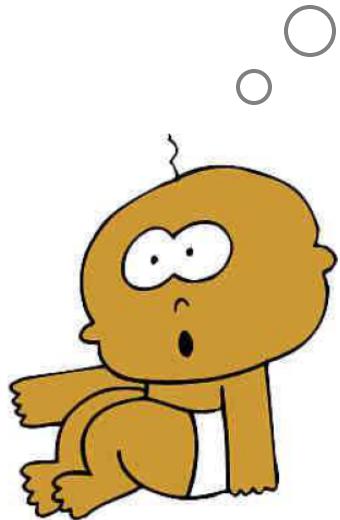
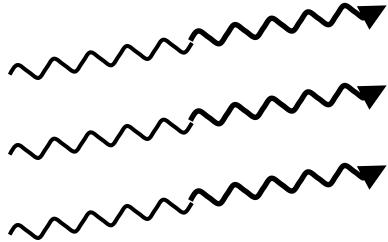
Curve fitting



Light



Sound



Language

Causes

Objects

Number

Folk Physics

Take home messages

- Bayesian models are useful for thinking about inductive problems.
- There are many inductive problems from a wide range of domains including vision, language, motor control, etc.

Module 2

BAYESIAN CONCEPT LEARNING

Learning from positive examples



“wombat”

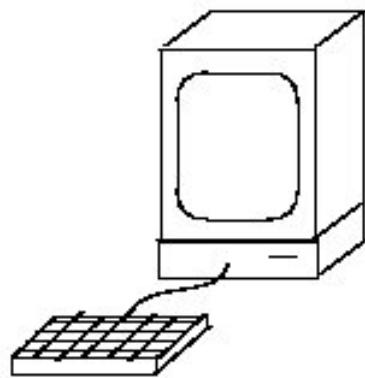


wombat?



wombat?

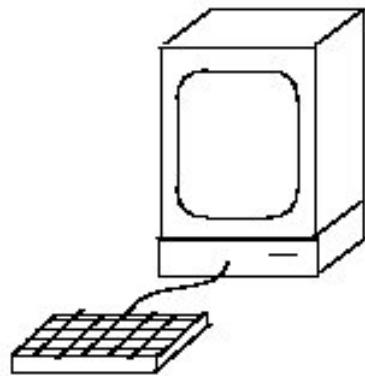
Learning from positive examples



30 70 40 50

(Tenenbaum)

Learning from positive examples



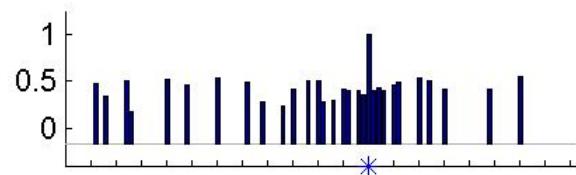
72 68 73

(Tenenbaum)

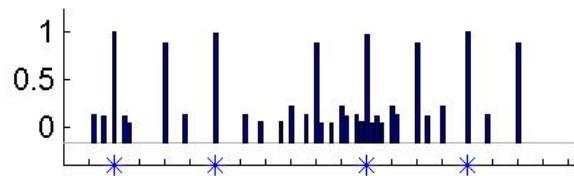
+ Examples

Human generalization Bayesian Model

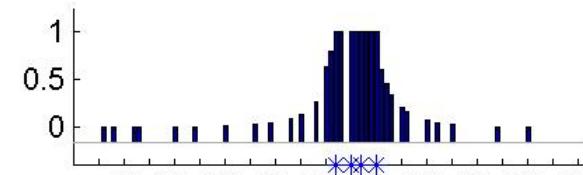
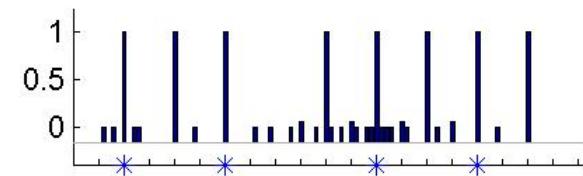
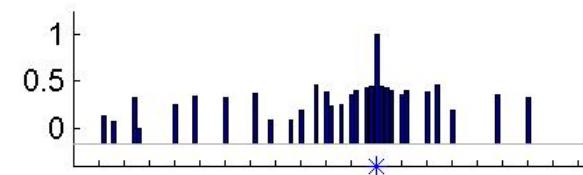
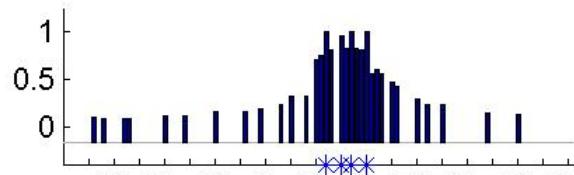
60



60 80 10 30



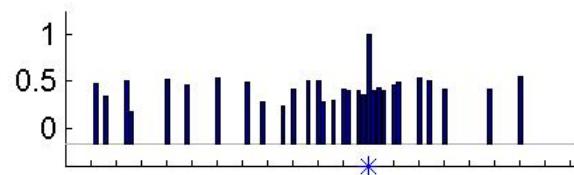
60 52 57 55



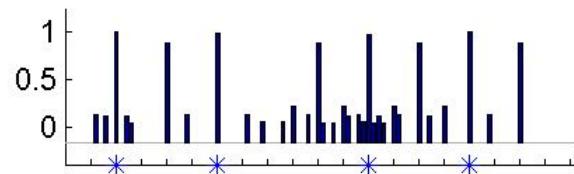
+ Examples

Human generalization Bayesian Model

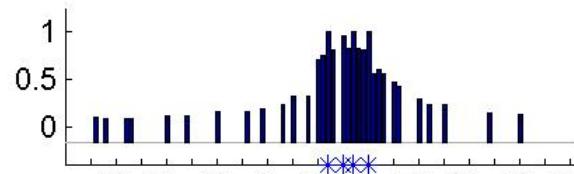
60



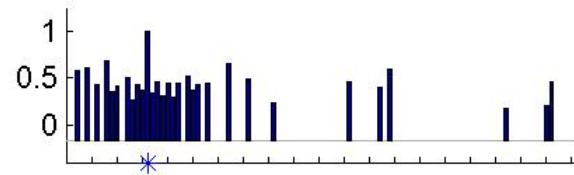
60 80 10 30



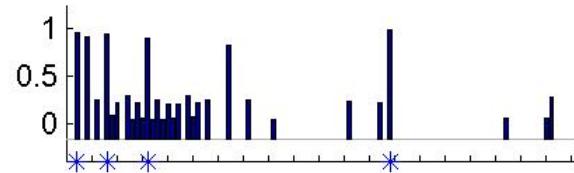
60 52 57 55



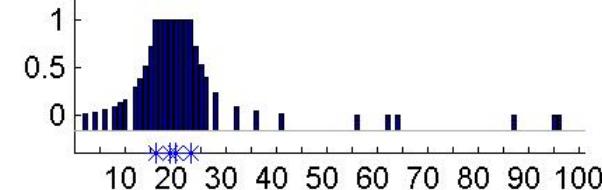
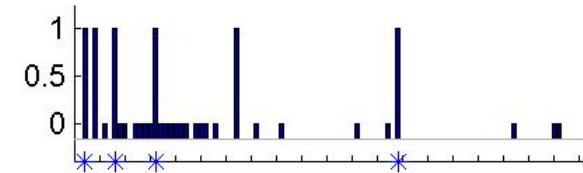
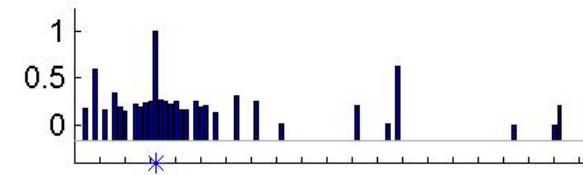
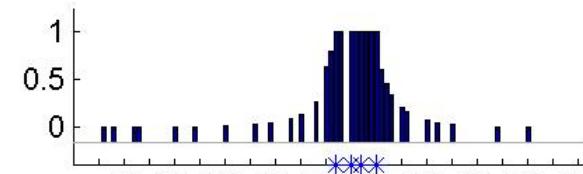
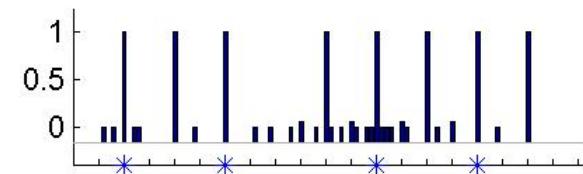
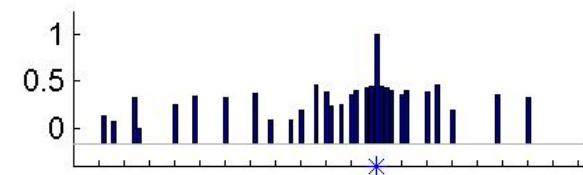
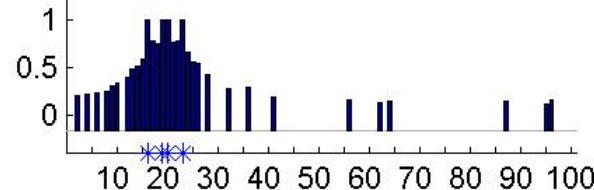
16



16 8 2 64



16 23 19 20



Building a Bayesian model

- What are the observed data?
- What are the hypotheses?
- What is the prior?
- What is the likelihood?

Posterior

Likelihood

Prior



$$P(h|d) \propto P(d|h) P(h)$$

Bayesian model

- H : Hypothesis space of possible concepts.
 - $h_1 = \{2, 4, 6, 8, 10, 12, \dots, 96, 98, 100\}$ (“even numbers”)
 - $h_2 = \{10, 20, 30, 40, \dots, 90, 100\}$ (“multiples of 10”)
 - $h_3 = \{2, 4, 8, 16, 32, 64\}$ (“powers of 2”)
 - $h_4 = \{50, 51, 52, \dots, 59, 60\}$ (“numbers between 50 and 60”)
- $X = \{x_1, \dots, x_n\}$: n examples of a concept C .
- Posterior probability of hypothesis h given the data

$$P(h|X) = \frac{P(X|h) P(h)}{P(X)}$$

Prior $P(h)$

- Mathematical properties (24 hypotheses):
 - Odd, even, square, cube, prime numbers
 - Multiples of small integers
 - Powers of small integers
- Raw magnitude (5050 hypotheses):
 - All intervals of integers with endpoints between 1 and 100.

Likelihood $P(X|h)$ (strong sampling)

- Assume that the examples are drawn at random from the hypothesis:

$$P(X|h) = \begin{cases} \left[\frac{1}{\text{size}(h)} \right]^n & \text{if all } x_i \text{ are in } h \\ 0 & \text{otherwise} \end{cases}$$

Illustrating the likelihood

2	4	6	8	10
12	14	16	18	20
22	24	26	28	30
32	34	36	38	40
42	44	46	48	50
52	54	56	58	60
62	64	66	68	70
72	74	76	78	80
82	84	86	88	90
92	94	96	98	100

h_1 →

← h_2

Illustrating the likelihood

2	4	6	8	10
12	14	16	18	20
22	24	26	28	30
32	34	36	38	40
42	44	46	48	50
52	54	56	58	60
62	64	66	68	70
72	74	76	78	80
82	84	86	88	90
92	94	96	98	100

Data slightly more of a coincidence under h_1

Illustrating the likelihood

2	4	6	8	10
12	14	16	18	20
22	24	26	28	30
32	34	36	38	40
42	44	46	48	50
52	54	56	58	60
62	64	66	68	70
72	74	76	78	80
82	84	86	88	90
92	94	96	98	100

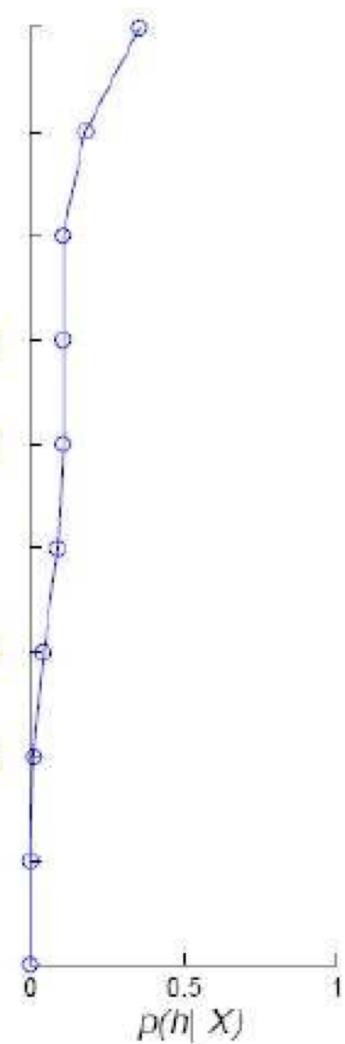
h_1 →

← h_2

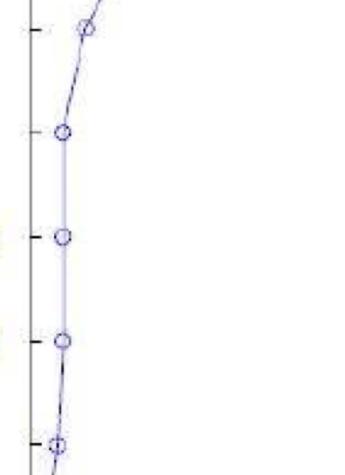
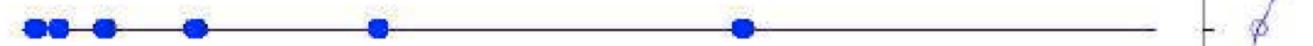
Data *much* more of a coincidence under h_1

Examples: 16

powers of 4



powers of 2



nos. ending in 6



square numbers



even numbers



multiples of 8



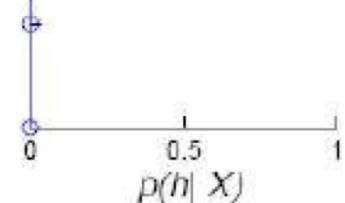
multiples of 4



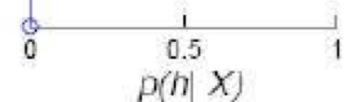
nos. 1–100



powers of 2, - 32



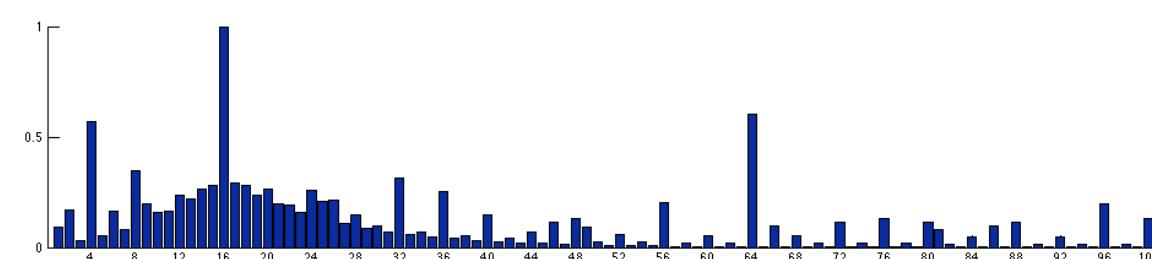
powers of 2, + 37



Prediction by hypothesis averaging

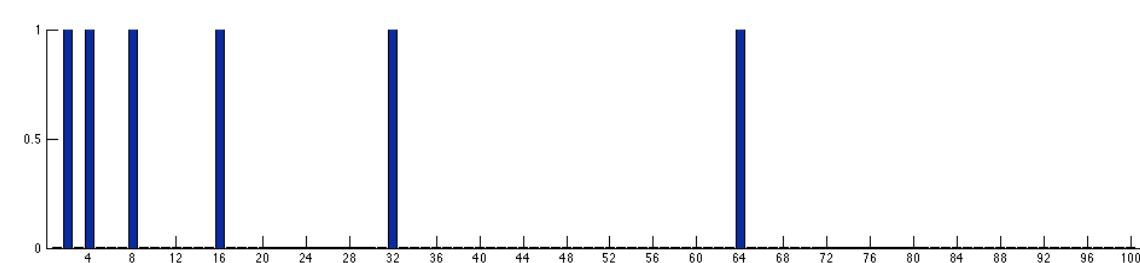
$$\begin{aligned} P(y \in C|X) &= \sum_{h \in \mathcal{H}} P(y \in C|h)P(h|X) \\ &= \sum_{h \in \mathcal{H}_y} P(h|X) \end{aligned}$$

Tenenbaum prior,
strong sampling



16 8 2 64

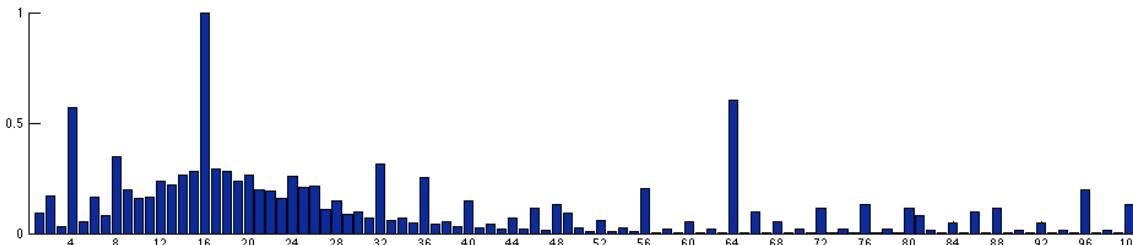
Tenenbaum prior,
strong sampling



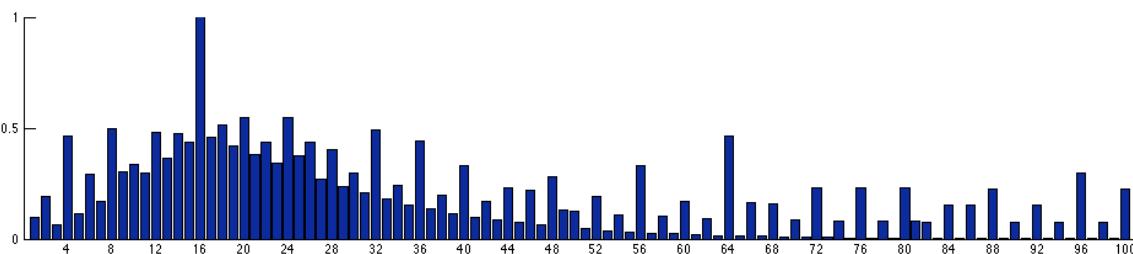
Likelihood $P(X|h)$ (weak sampling)

- Assume that someone has asked whether each x_i is in the concept

$$\begin{aligned} P(X|h) &= 1 && \text{if all } x_i \text{ are in } h \\ &= 0 && \text{otherwise} \end{aligned}$$



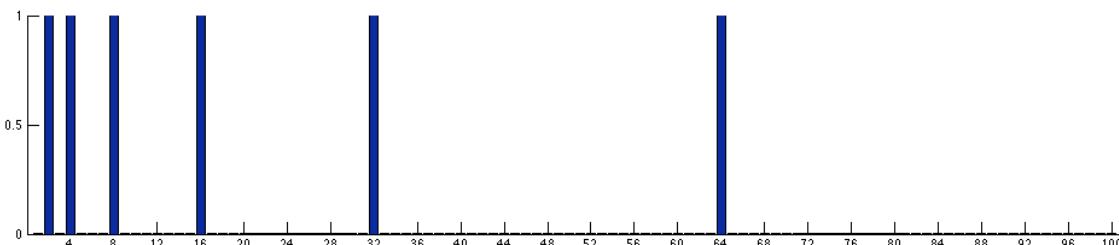
Tenenbaum prior,
strong sampling



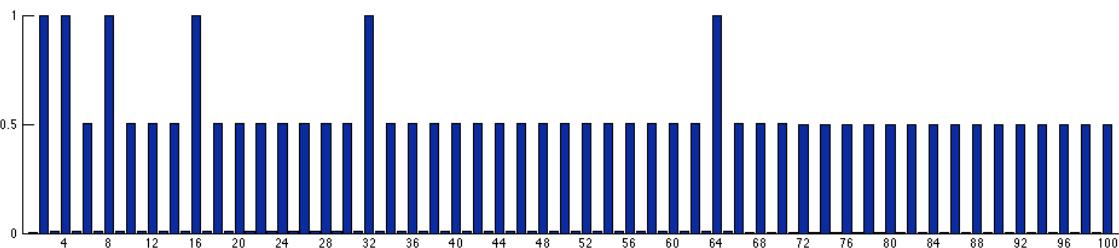
Tenenbaum prior,
weak sampling

16 8 2 64

Tenenbaum prior,
strong sampling



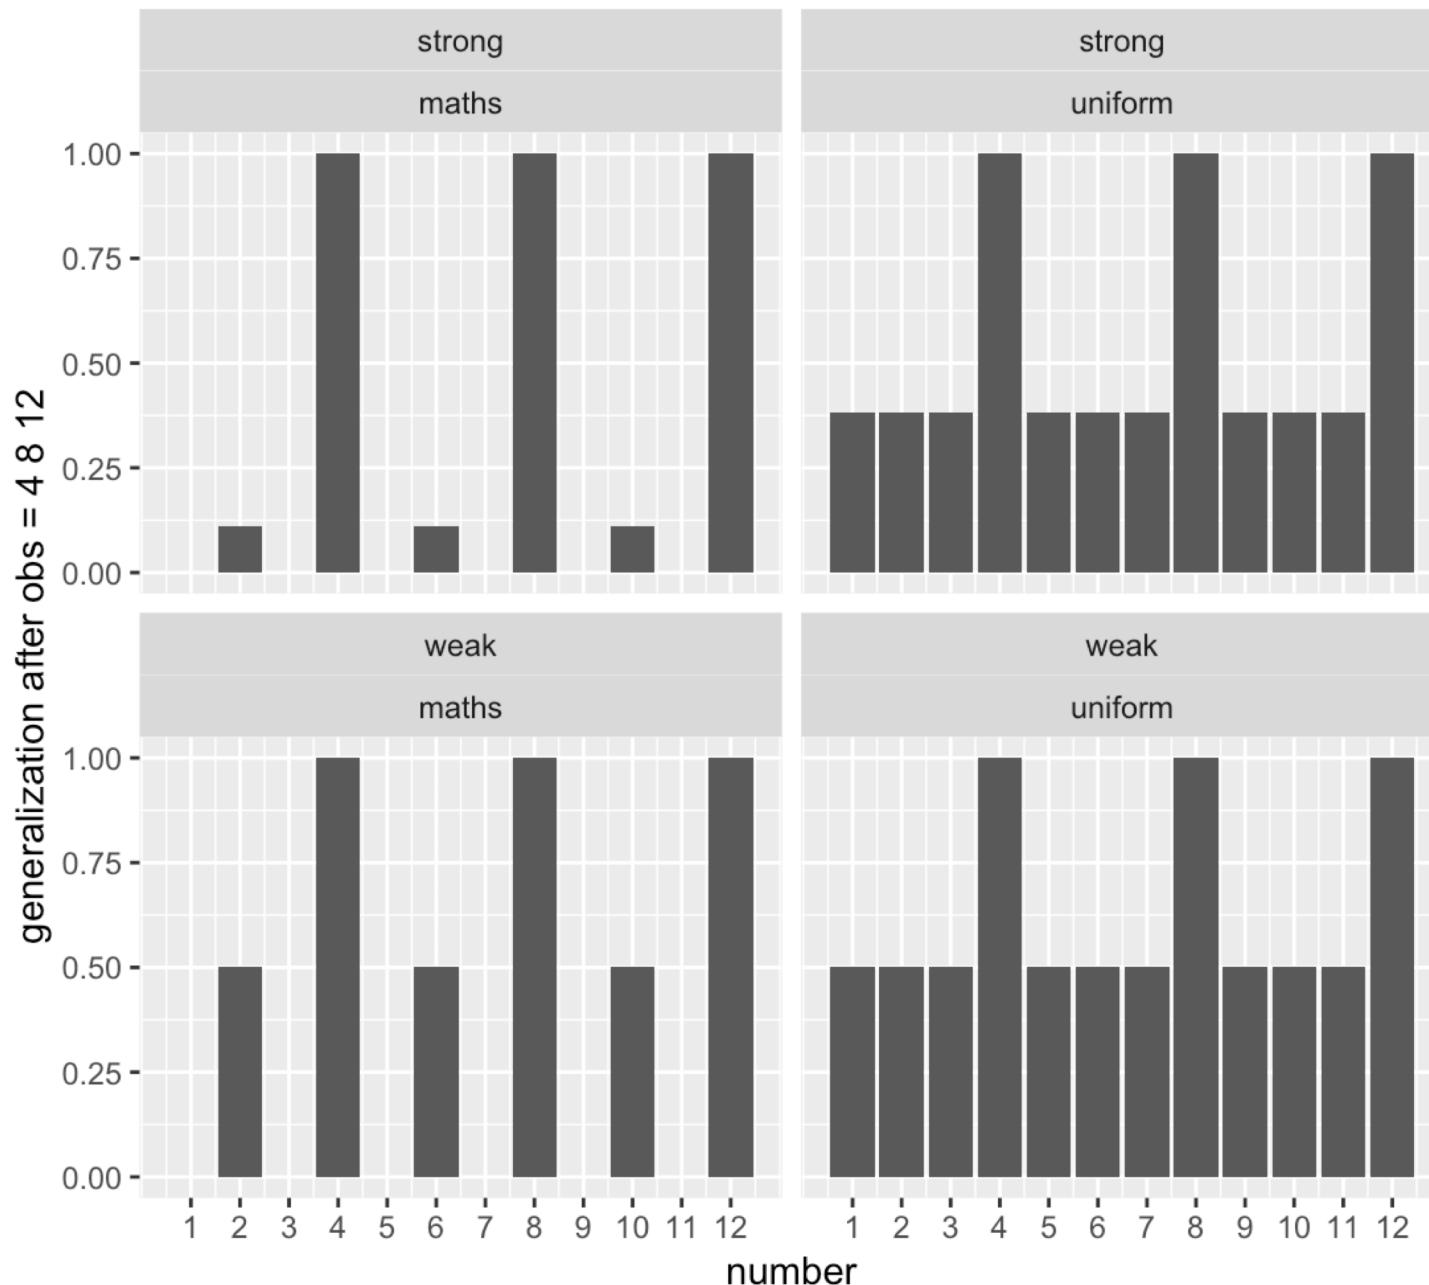
Tenenbaum prior,
weak sampling



Take home messages

- Bayesian models include a likelihood and a prior.
- The likelihood matters.
- How much does the prior matter?

Exercise: number game



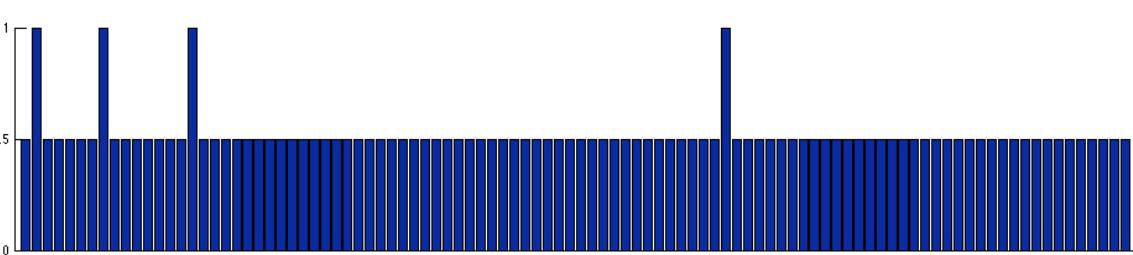
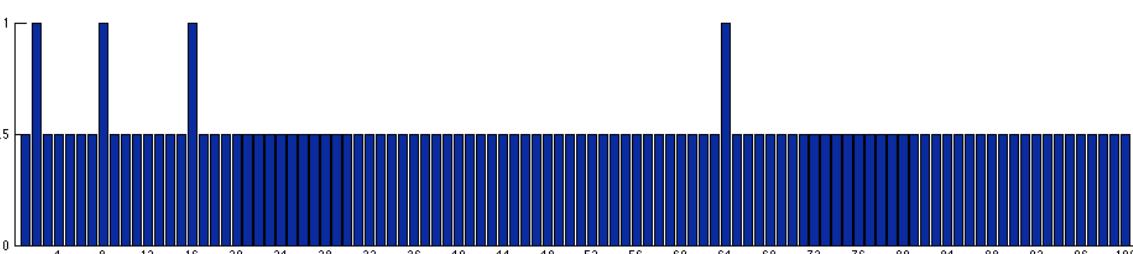
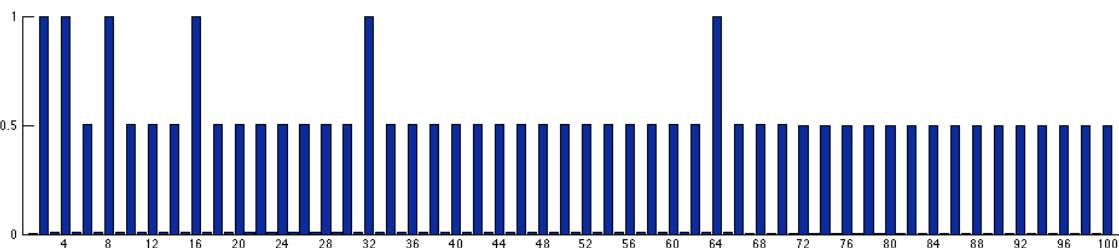
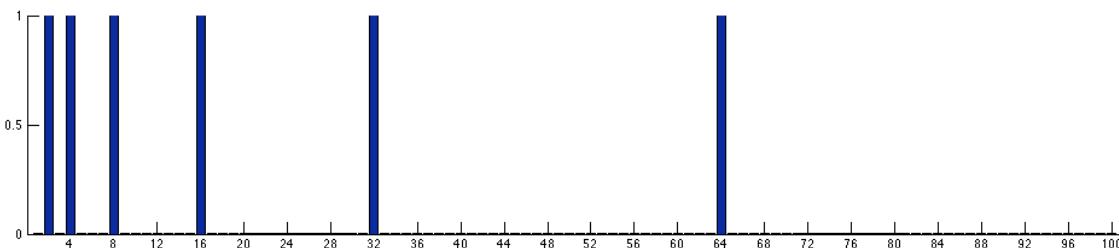
16 8 2 64

Tenenbaum prior,
strong sampling

Tenenbaum prior,
weak sampling

Uniform prior,
strong sampling

Uniform prior,
weak sampling



Take home messages

- Bayesian models include a likelihood and a prior.
- The likelihood matters.
- The prior matters.
- Think carefully about both when you set up a model!

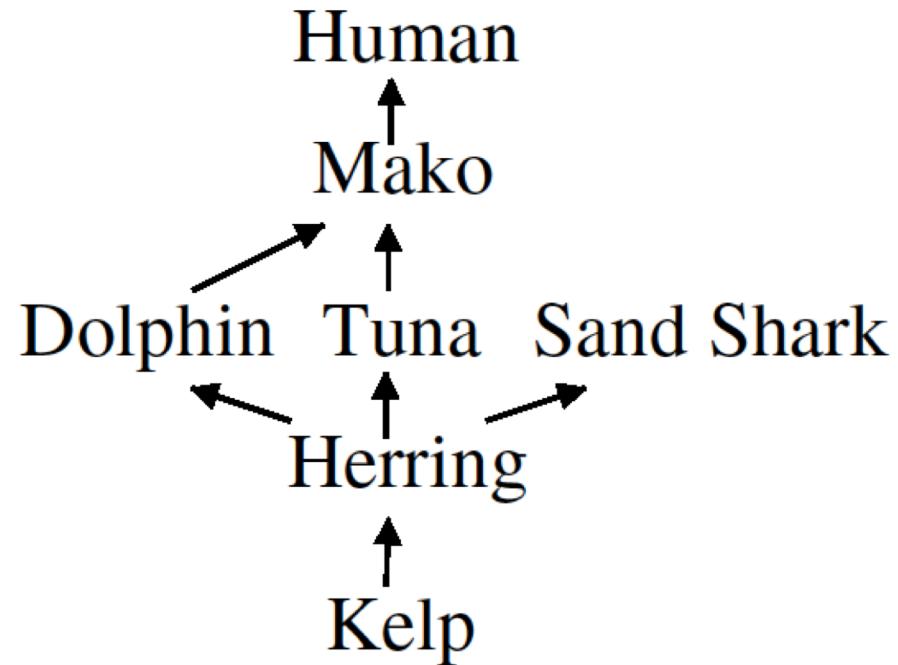
Module 3

BAYESIAN NETWORKS (DIRECTED GRAPHICAL MODELS)

Representing hypothesis spaces and priors

- So far we've just enumerated the hypothesis space.
- What if the hypothesis space is huge?
- How do we come up with all the numbers in the prior?

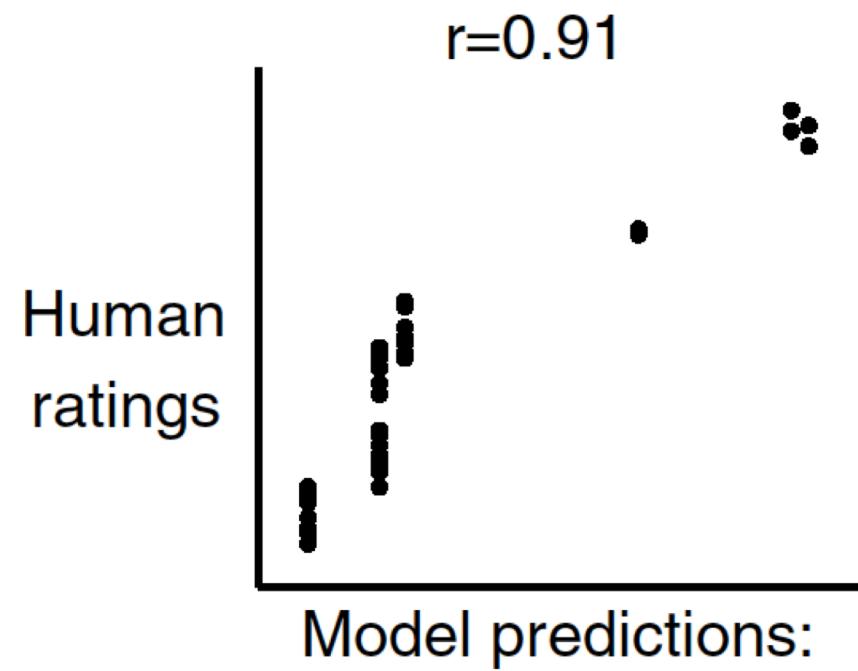
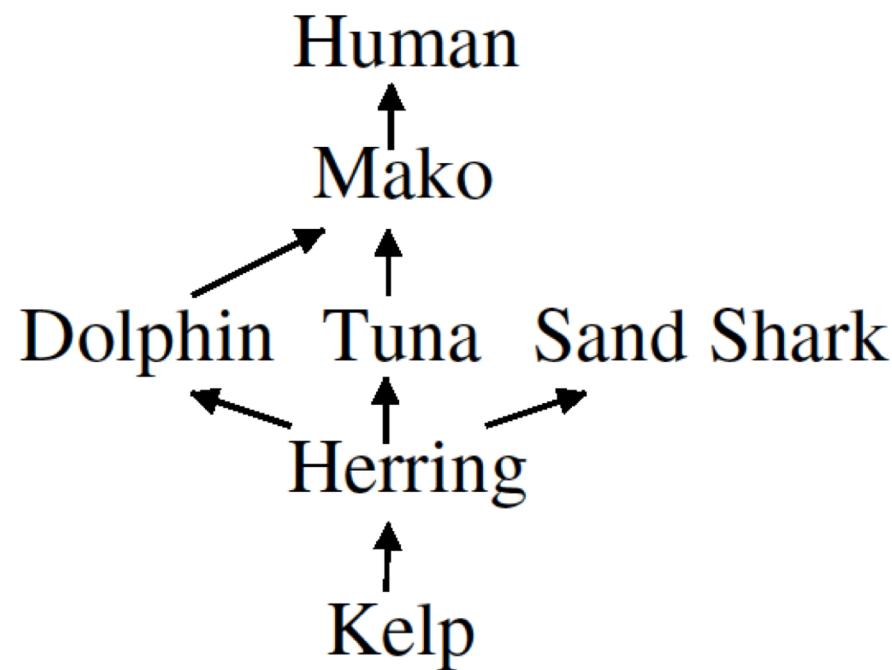
Food web problem



Herring carry a certain disease.

How likely is it that mako have the same disease?

Food web problem



(Shafto et al, 2008)

Building a Bayesian model

- What are the observed data?
- What are the hypotheses?
- What is the prior?
- What is the likelihood?

Posterior

Likelihood

Prior

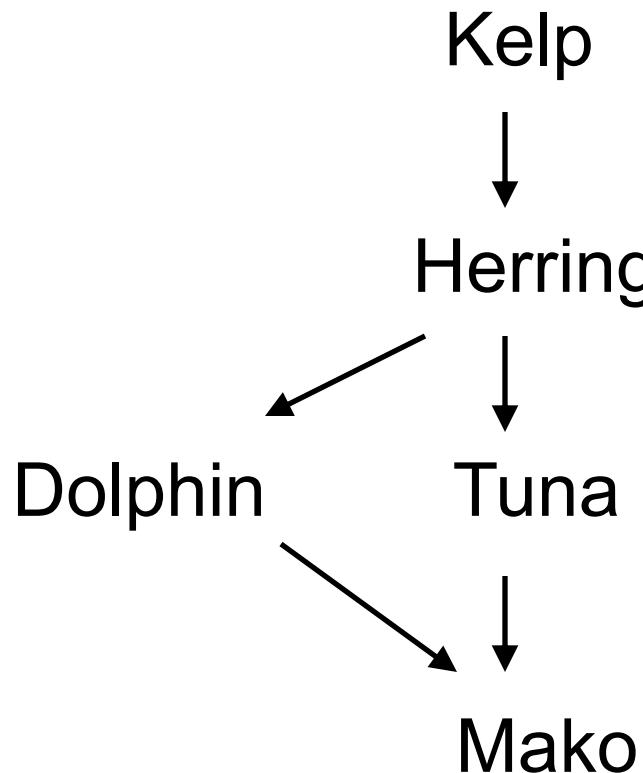


$$P(h|d) \propto P(d|h) P(h)$$

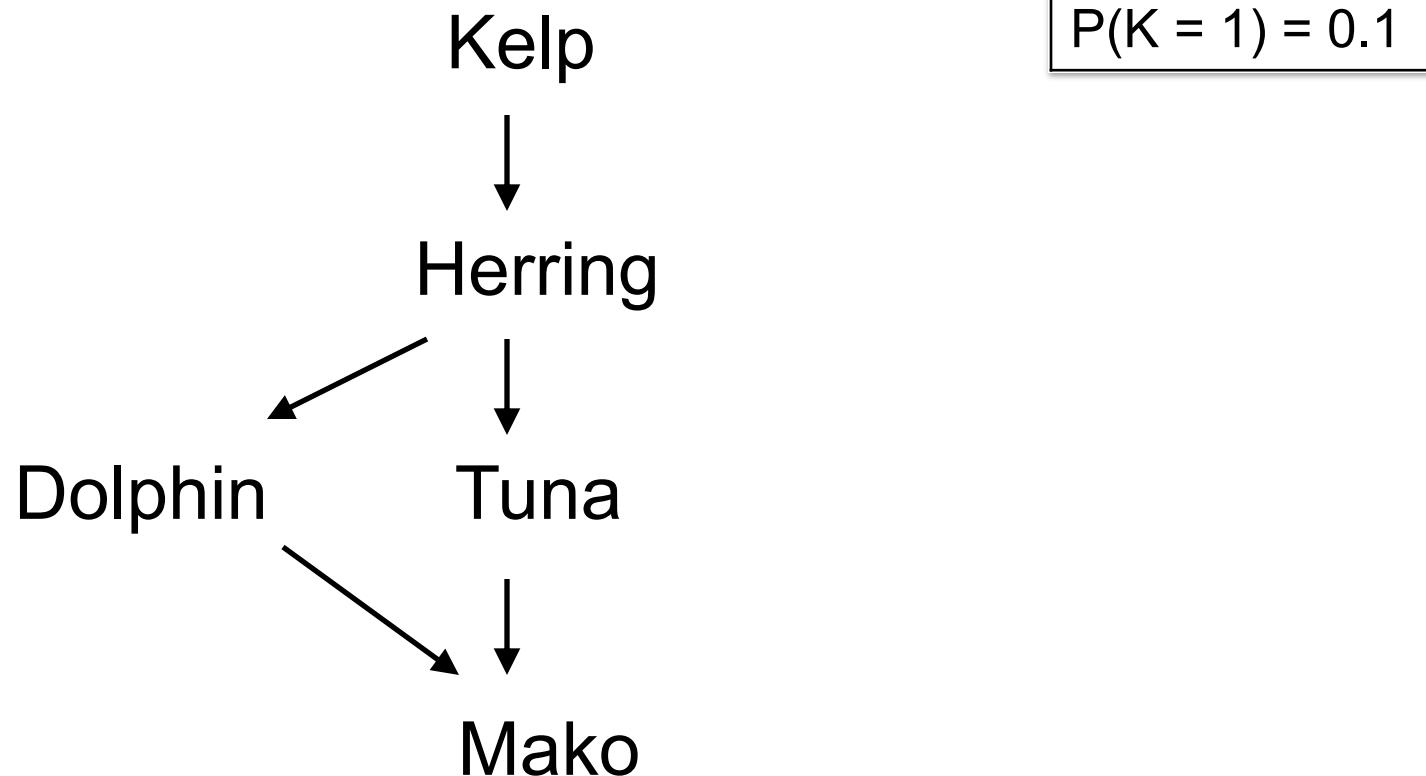
Hypothesis space and prior

	kelp	herring	dolphin	tuna	sandshark	mako	human	prior
	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>
1	1	1	1	1	1	1	1	0.478
2	2	1	1	1	1	1	1	0.0266
3	1	2	1	1	1	1	1	0.00664
4	2	2	1	1	1	1	1	0.00406
5	1	1	2	1	1	1	1	0.0266
6	2	1	2	1	1	1	1	0.00148
7	1	2	2	1	1	1	1	0.00406
8	2	2	2	1	1	1	1	0.00248
9	1	1	1	2	1	1	1	0.0266
10	2	1	1	2	1	1	1	0.00148

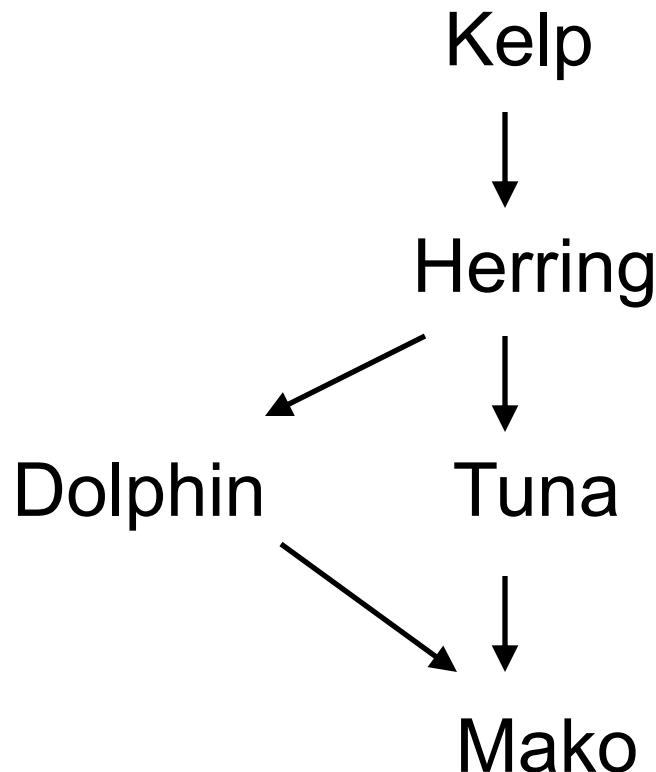
Specifying a prior



Specifying a prior



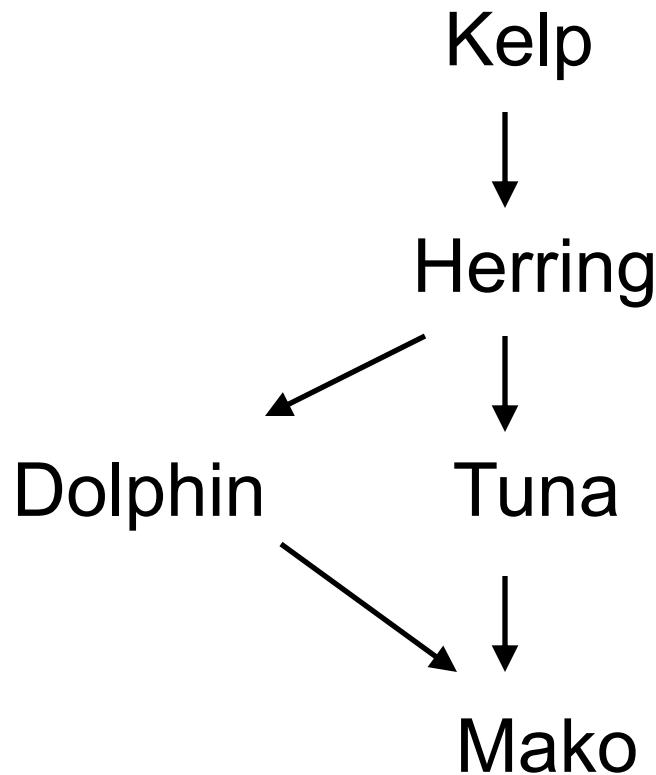
Specifying a prior



$$P(K = 1) = 0.1$$

K	$P(H = 1 K)$
0	0.1
1	0.55

Specifying a prior

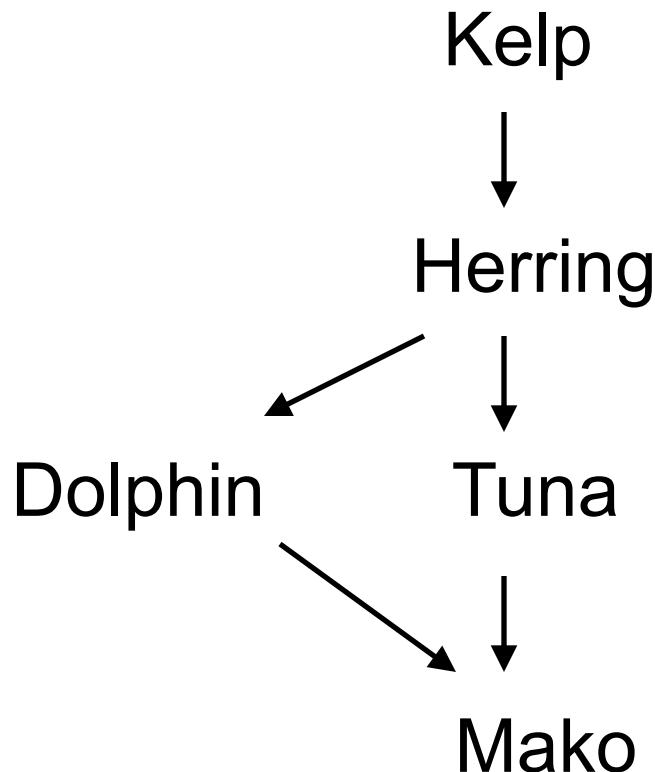


$$P(K = 1) = 0.1$$

K	P(H = 1 K)
0	0.1
1	0.55

D	T	P(M = 1 D, T)
0	0	0.1
0	1	0.55
1	0	0.55
1	1	0.775

Specifying a prior

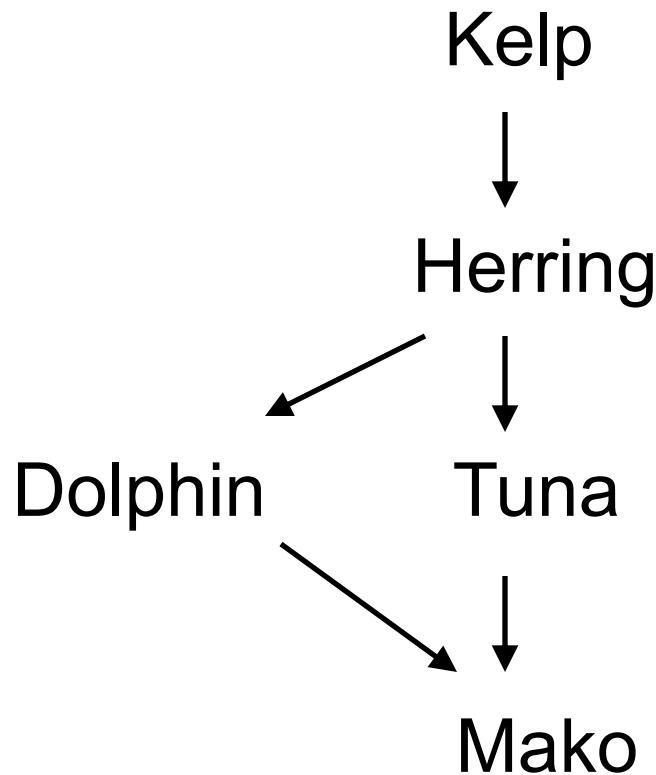


$$\begin{aligned} P(K, H, D, T, M) = \\ P(K) * \\ P(H|K) * \\ P(D|H) * \\ P(T|H) * \\ P(M|D, T) \end{aligned}$$

General formulation

$$P(v_1, \dots, v_n) = \prod_i P(v_i | \text{pa}(V_i))$$

Specifying a prior



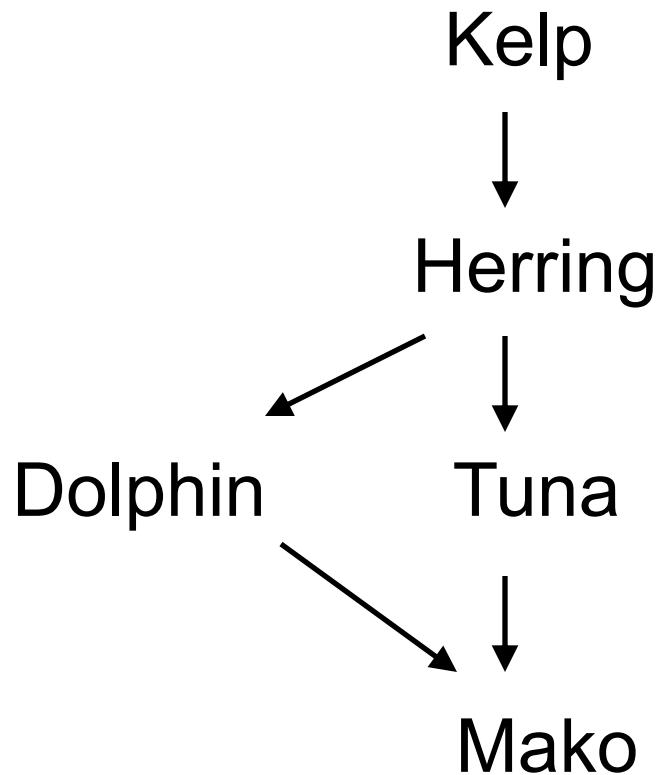
$$P(K = 1) = 0.1$$

K	P(H = 1 K)
0	0.1
1	0.55

D	T	P(M = 1 D, T)
0	0	0.1
0	1	0.55
1	0	0.55
1	1	0.775

Specifying a prior

b: base rate
t: transmission rate



$$P(K = 1) = b$$

K	$P(H = 1 K)$
0	b
1	$1 - (1-t)(1-b)$

D	T	$P(M = 1 D, T)$
0	0	b
0	1	$1 - (1-t)(1-b)$
1	0	$1 - (1-t)(1-b)$
1	1	$1 - (1-t)(1-t)(1-b)$

Exercise: Food web (enumeration)

	kelp	herring	dolphin	tuna	sandshark	mako	human	prior
	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>
1	1	1	1	1	1	1	1	0.478
2	2	1	1	1	1	1	1	0.0266
3	1	2	1	1	1	1	1	0.00664
4	2	2	1	1	1	1	1	0.00406
5	1	1	2	1	1	1	1	0.0266
6	2	1	2	1	1	1	1	0.00148
7	1	2	2	1	1	1	1	0.00406
8	2	2	2	1	1	1	1	0.00248
9	1	1	1	2	1	1	1	0.0266
10	2	1	1	2	1	1	1	0.00148

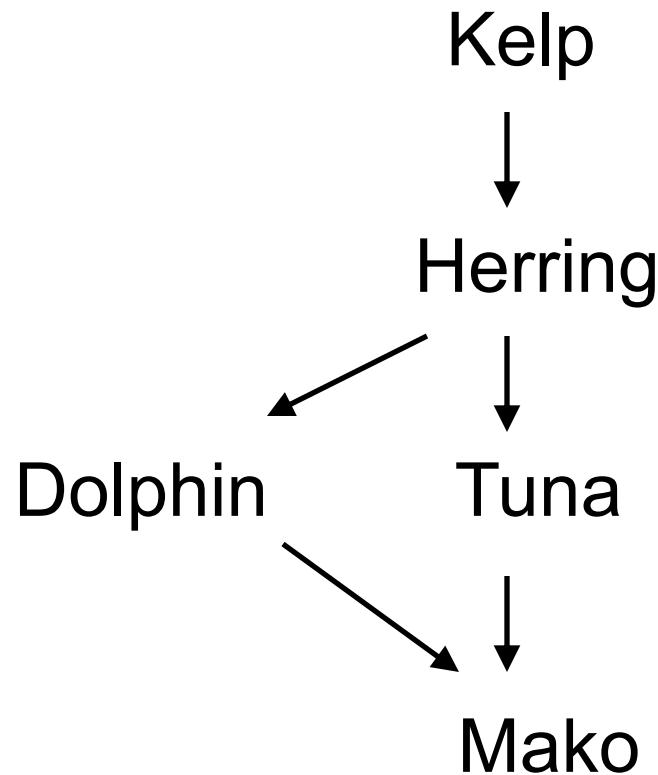
Take home messages

- Bayes nets give us a good way to specify priors

Module 4

INFERENCE BY SAMPLING FROM THE PRIOR

Sampling from the prior



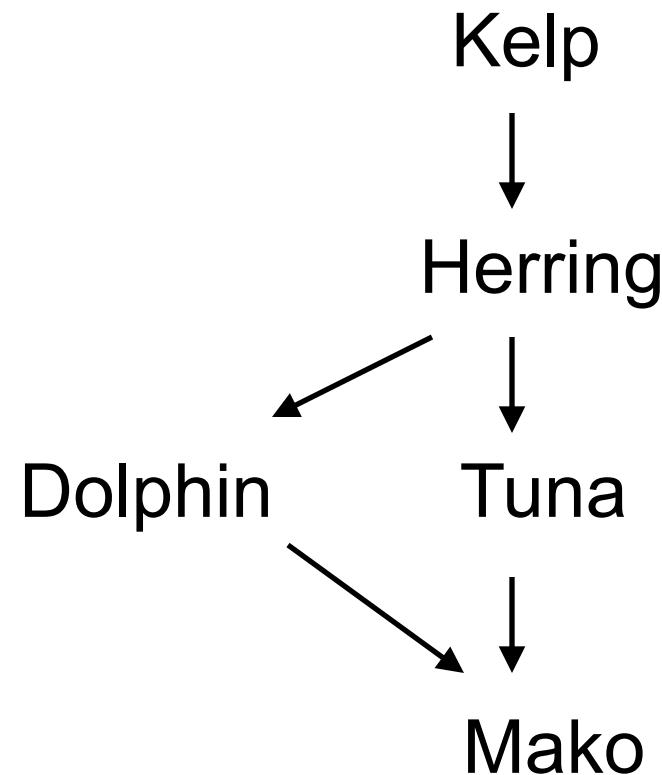
$$P(K = 1) = 0.1$$

K	P(H = 1 K)
0	0.1
1	0.55

D	T	P(M = 1 D, T)
0	0	0.1
0	1	0.55
1	0	0.55
1	1	0.775

Sample based predictions

```
obs <- list(kelp = 1, mako = 2)
```



1. Collect set of samples from prior
2. Remove all those that are inconsistent with the observations
3. Make predictions based on samples that remain

Sample based predictions

$$\begin{aligned} P(\text{humans}|obs) &= \sum_h P(\text{humans}|h)P(h|obs) \\ &\propto \sum_h P(\text{humans}|h)P(obs|h)P(h) \\ &\approx \frac{1}{M} \sum_{i=1}^M P(\text{humans}|h^i)P(obs|h^i) \end{aligned}$$

Exercise: Food web (sampling from the prior)

Take home messages

- Bayes nets give us a good way to specify priors
- Sampling is often a good way to implement probabilistic inference

Module 5

JAGS

Sample based predictions

- If we could sample from the posterior $P(h|obs)$

$$P(\text{humans}|obs) = \sum_h P(\text{humans}|h)P(h|obs)$$

$$\approx \frac{1}{M} \sum_{i=1}^M P(\text{humans}|h^i)$$

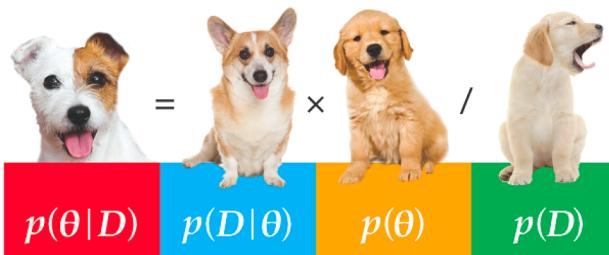
JAGS and STAN

- Widely used for data analysis and cognitive modeling

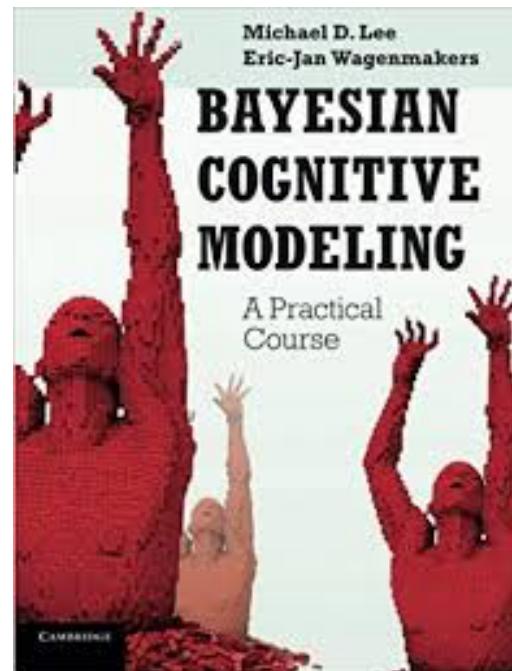
Second Edition

Doing Bayesian Data Analysis

A Tutorial with R, JAGS, and Stan



John K. Kruschke



JAGS : model specification

```
model
{
  # dcat specifies a discrete categorical distribution
  kelp ~ dcat(p.kelp[1:2])
  herring ~ dcat(p.herring[kelp,1:2])
  dolphin ~ dcat(p.dolphin[herring,1:2])
  tuna ~ dcat(p.tuna[herring,1:2])
  sandshark ~ dcat(p.sandshark[herring,1:2])
  mako ~ dcat(p.mako[dolphin,tuna,1:2])
  human ~ dcat(p.human[mako,1:2])
}
```

JAGS : observations and CPDs

```
# specify observations here
obs <- list(kelp = 1, mako = 2)

foodwebdata<- c(obs, list(
  p.kelp = zerop,
  p.herring = onep,
  p.dolphin = onep,
  p.tuna = onep,
  p.sandshark = onep,
  p.mako = twop,
  p.human = onep))
```

Running JAGS

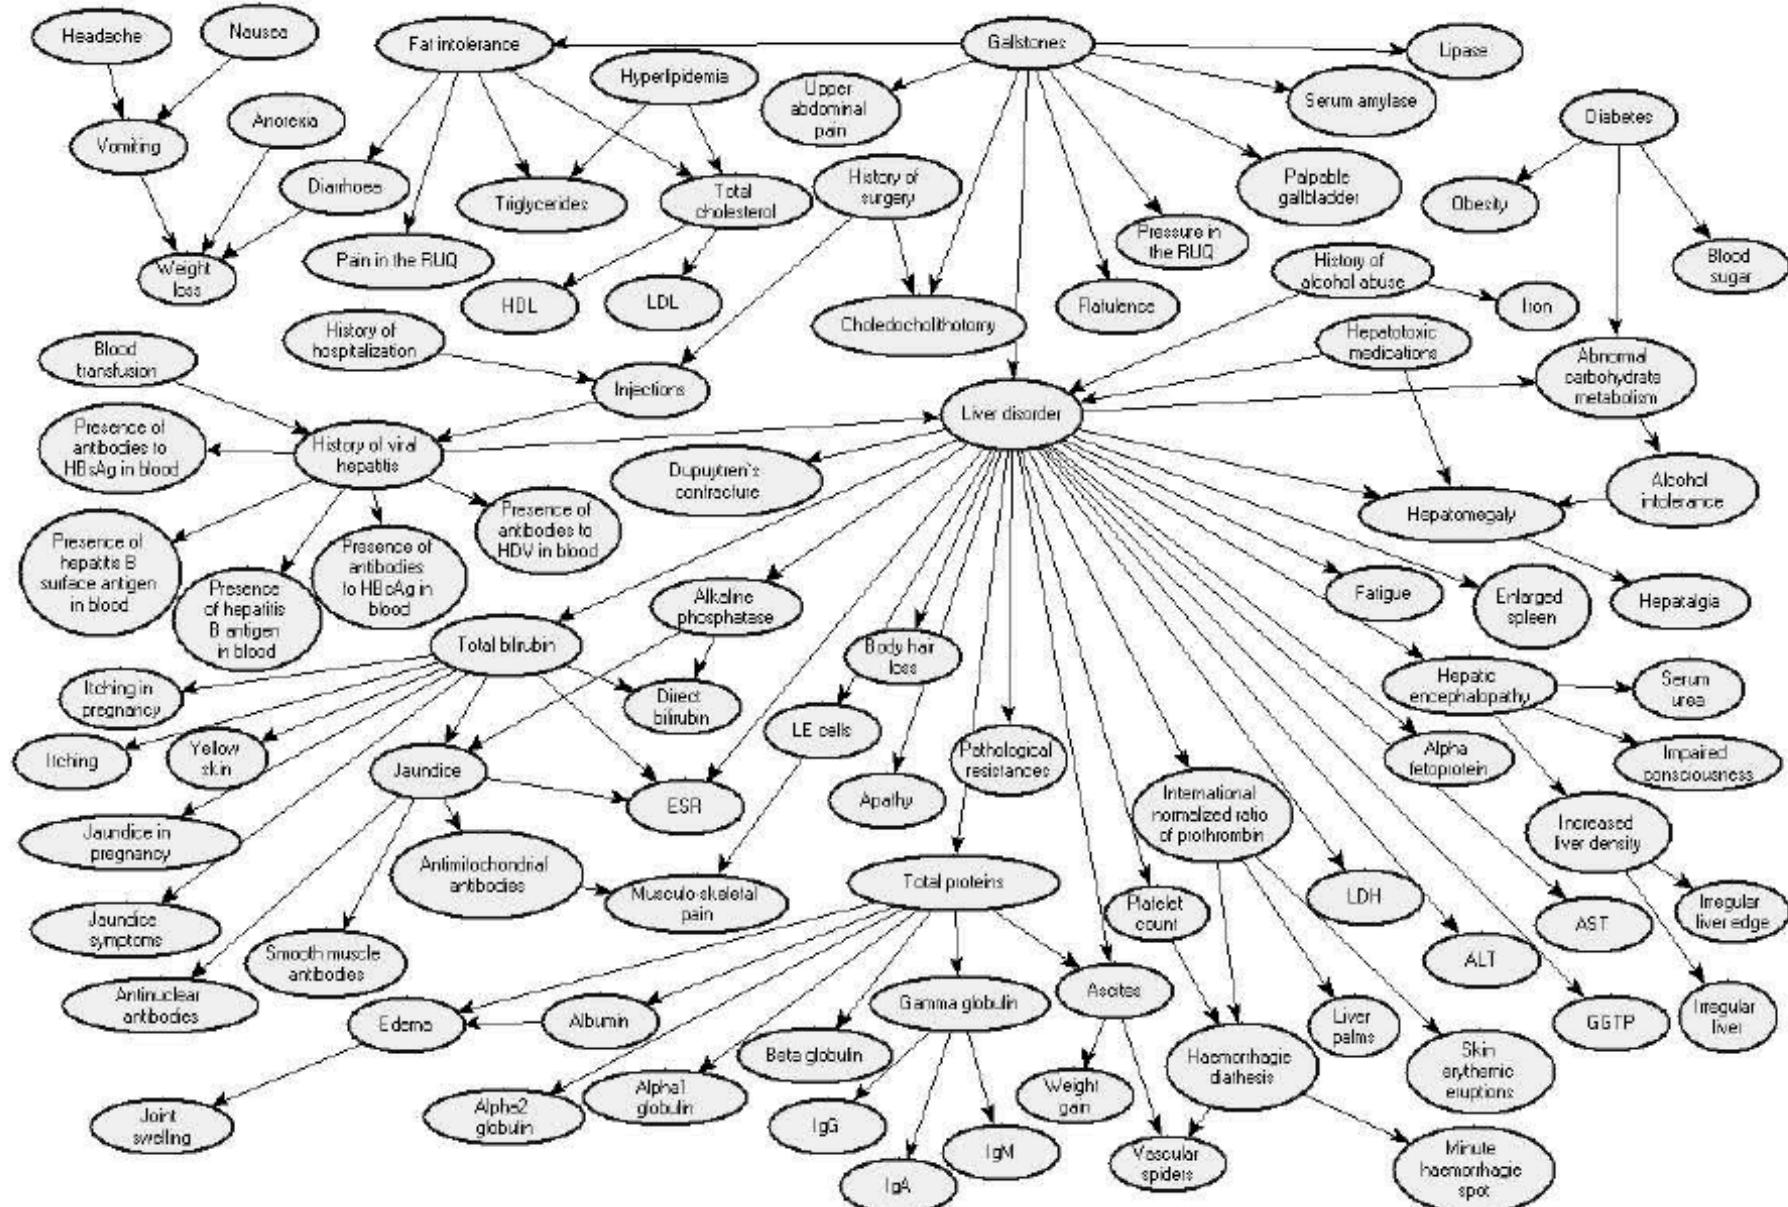
```
# set up the model in JAGS
jags <- jags.model('foodweb.bug', data = foodwebdata,
                    n.chains = 4,
                    n.adapt = 100)

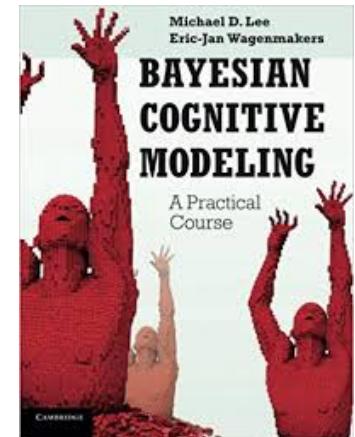
# actually run the model in JAGS (ie sample from the posterior P(h|obs) )
samples <- coda.samples(jags,
                        c('kelp', 'herring', 'dolphin', 'tuna', 'sandshark', 'mako', 'human'),
                        10000)
```

Exercise: Food web (sampling from the posterior using JAGS)

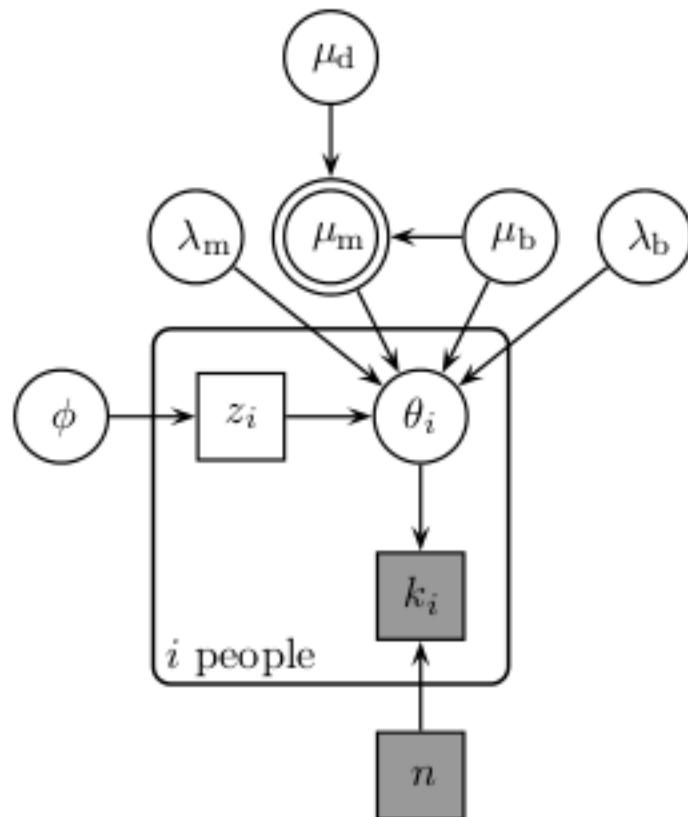
Take home messages

- Bayes nets give us a good way to specify priors
- Sampling is often a good way to implement probabilistic inference
- Tools like JAGS make sampling-based inference relatively simple in many contexts





Bayes nets



$$\mu_b \sim \text{Beta}(1, 1)$$

$$\mu_d \sim \text{Gaussian}(0, 0.5)_{\mathcal{I}(0, \infty)}$$

$$\lambda_b \sim \text{Uniform}(40, 800)$$

$$\lambda_m \sim \text{Uniform}(4, 100)$$

$$z_i \sim \text{Bernoulli}(\phi)$$

$$\theta_i \sim \begin{cases} \text{Beta}(\mu_b \lambda_b, (1 - \mu_b) \lambda_b) & \text{if } z_i = 0 \\ \text{Beta}(\mu_m \lambda_m, (1 - \mu_m) \lambda_m) & \text{if } z_i = 1 \end{cases}$$

$$k_i \sim \text{Binomial}(\theta_i, n)$$

$$\text{logit}\mu_m \leftarrow \text{logit}\mu_b - \mu_d$$

$$\phi \sim \text{Beta}(5, 5)$$

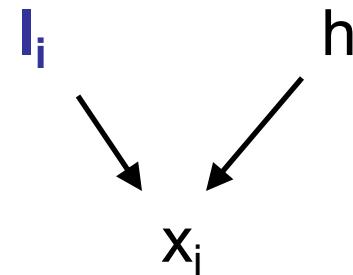
Graphical model for inferring membership of two latent groups, consisting of malingeringers and *bona fide* participants.

Bayes nets

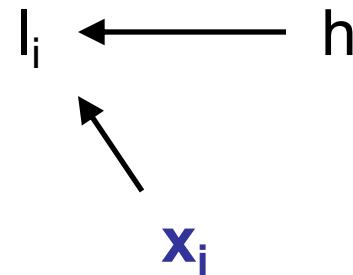
- When you're starting *any* modeling project
 - Try to write down the relevant variables
 - Draw a graph to show how the variables are related to each other

Bayes nets

- Strong sampling

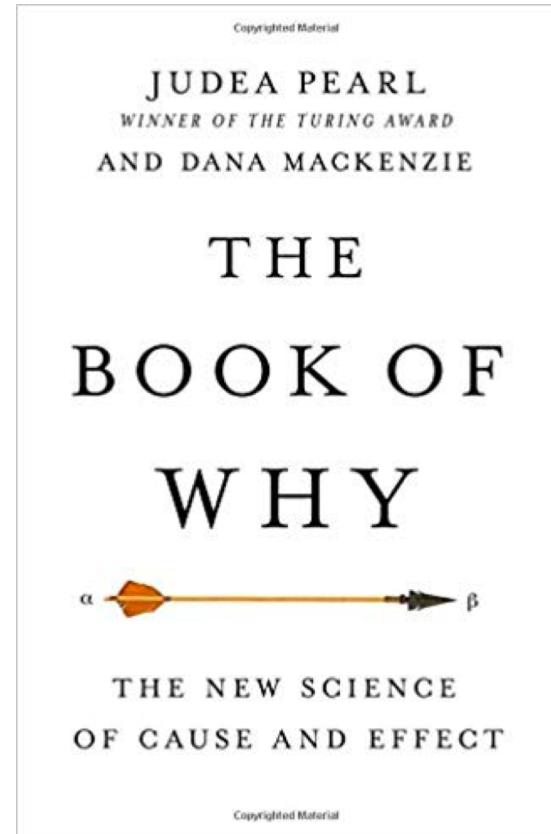


- Weak sampling



Bayes Nets and causal reasoning

- Bayes nets support reasoning about:
 - Interventions
 - Counterfactuals



Take home messages

- Bayes nets are:
 - a useful engineering tool
 - a tool for thinking

Module 6

SAMPLING FRAMES AND SPHERES OF SODOR

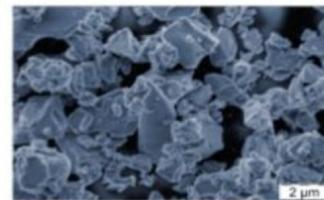
Sampling frames

The robot uses a grapple that can pick up small spheres and test them for the presence of plaxium



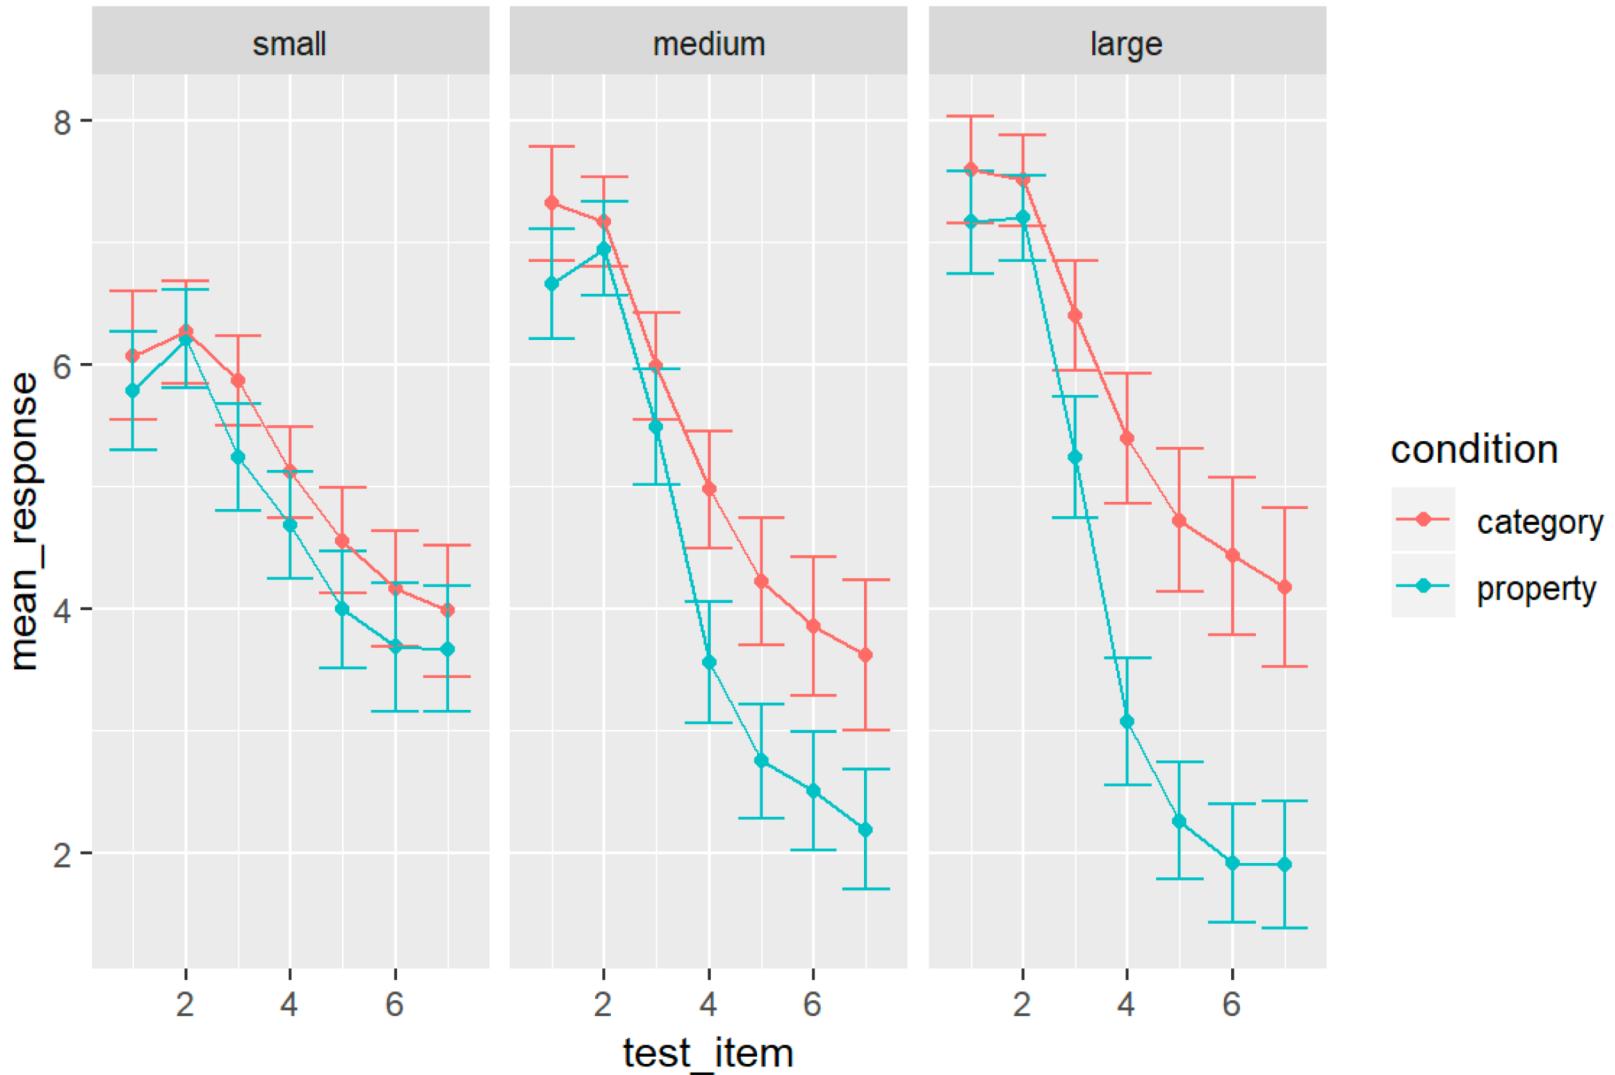
Category sampling

The robot uses a high-resolution scanner to detect plaxium, which triggers the camera to take an image whenever the plaxium coating is found



Property sampling

Human data



Building a Bayesian model

- What are the observed data?
- What are the hypotheses?
- What is the prior?
- What is the likelihood?

Posterior

Likelihood

Prior



$$P(h|d) \propto P(d|h) P(h)$$

Observations



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected



Plaxium detected

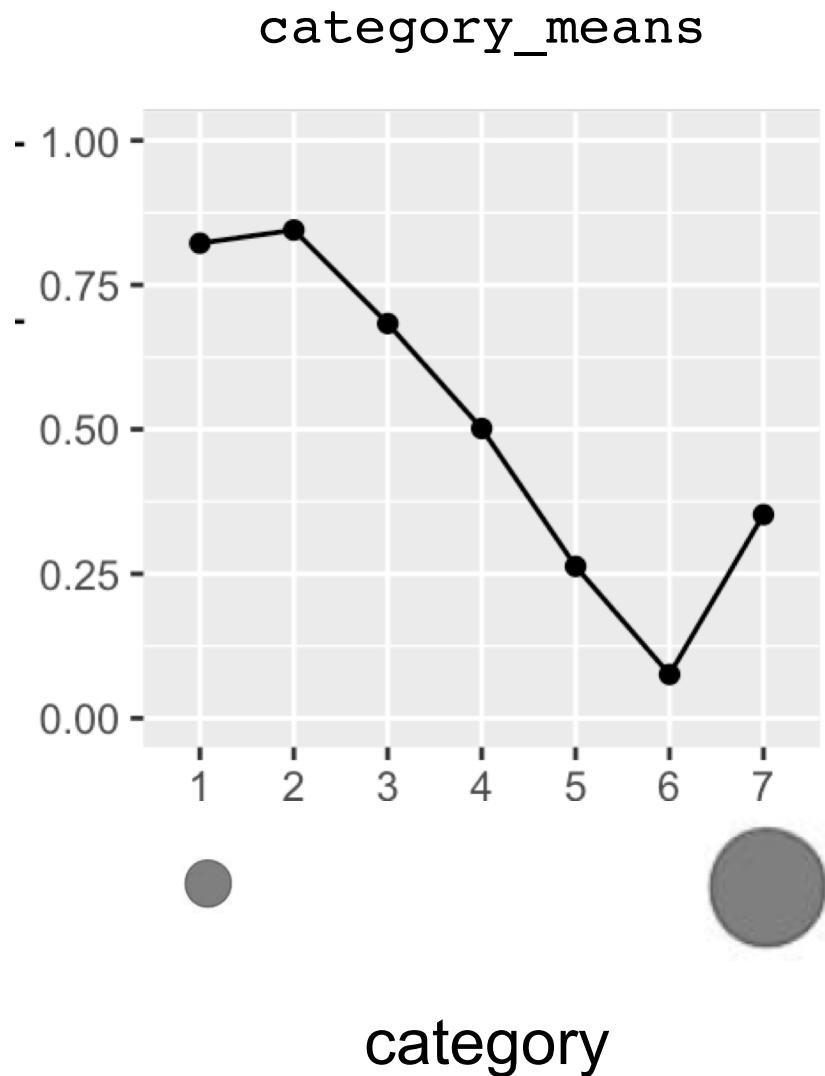


Plaxium detected

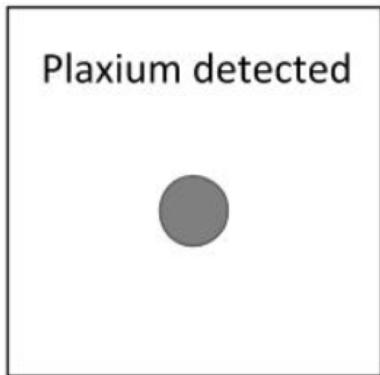


Hypotheses

Proportion
with
plaxium



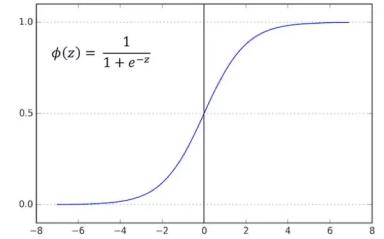
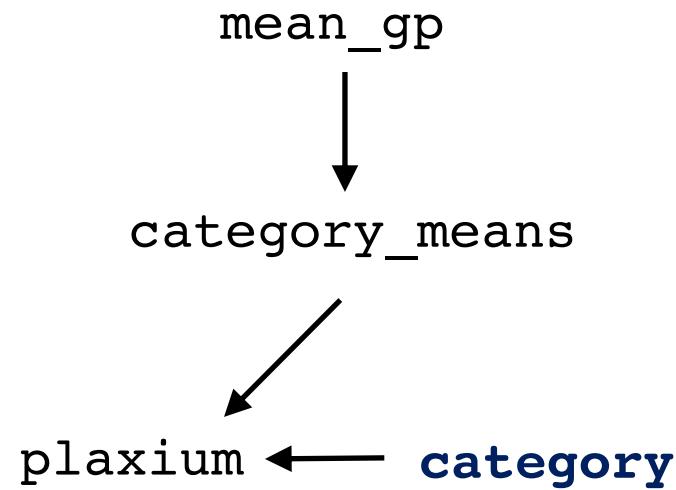
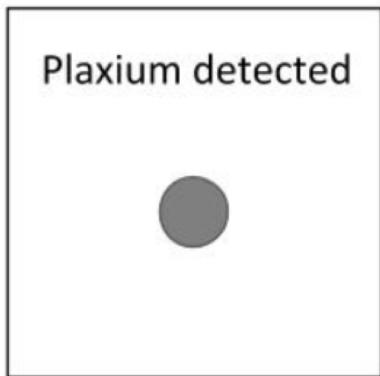
Category sampling



category_means

plaxium ← **category**

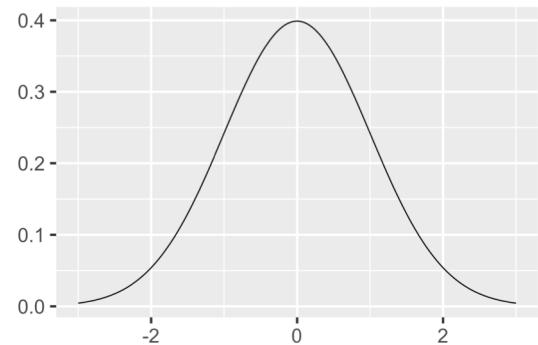
Category sampling



Gaussian distributions

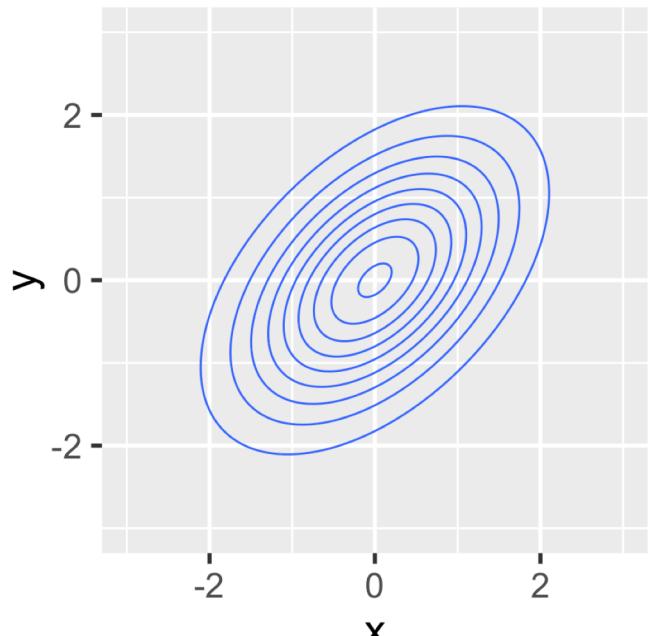
$$m = 0$$

$$\sigma = 1$$

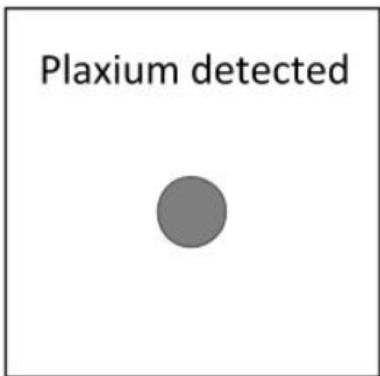
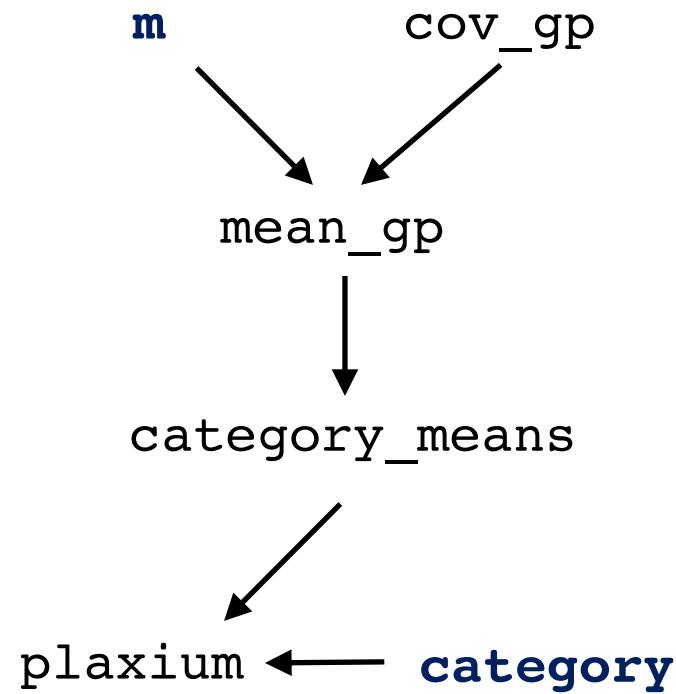


$$m = [0, 0]$$

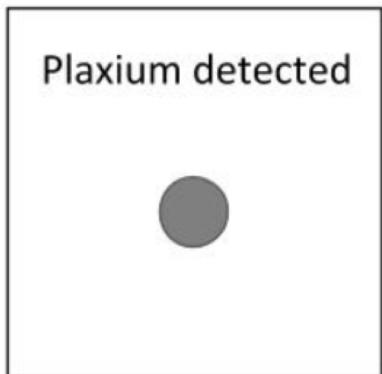
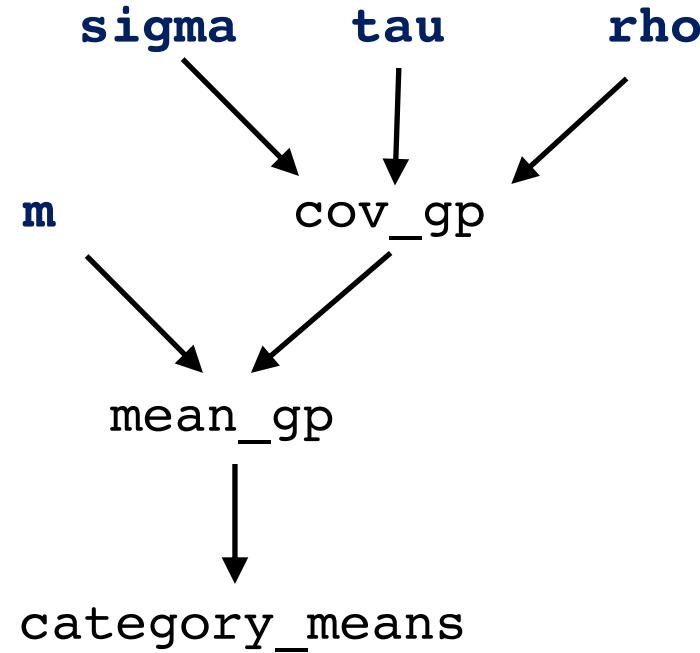
$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$$



Category sampling

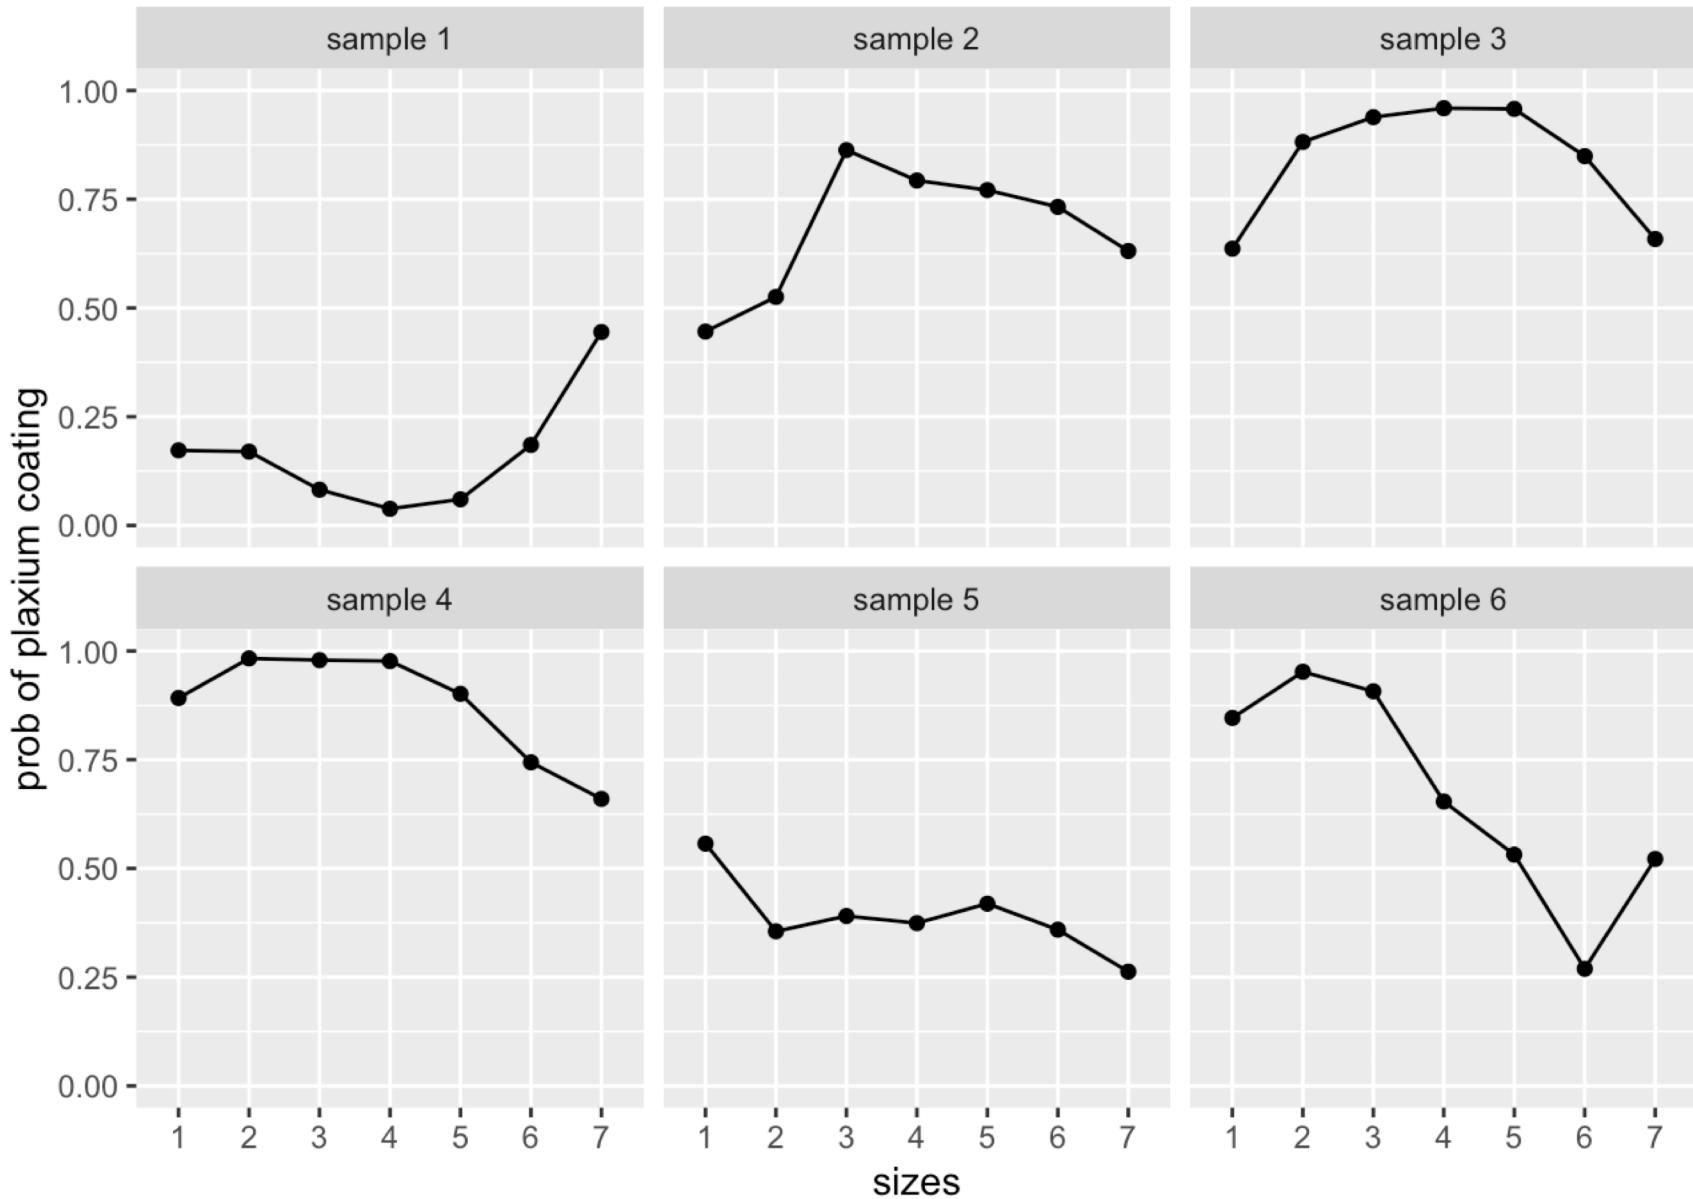


Category sampling

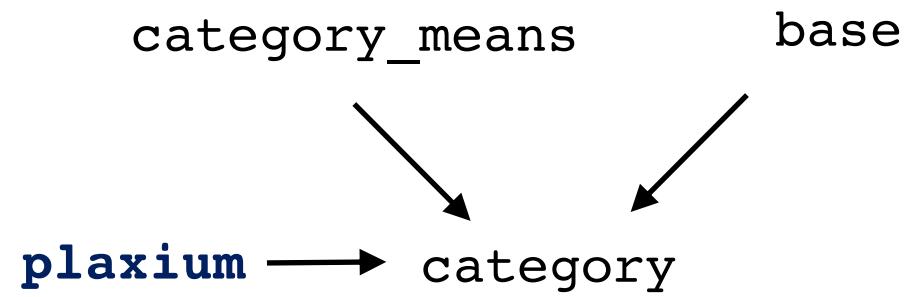
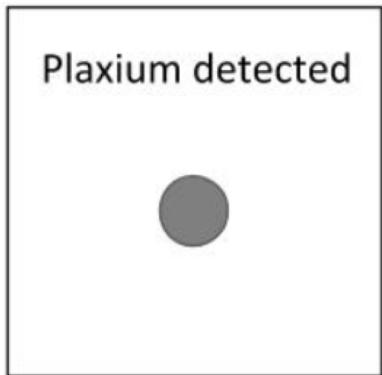


plaxium ← **category**

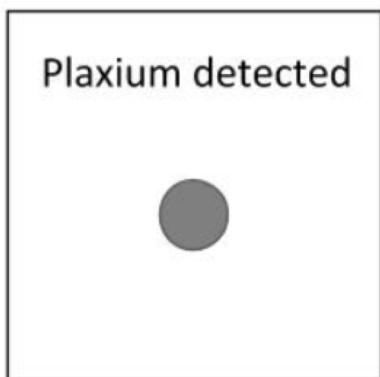
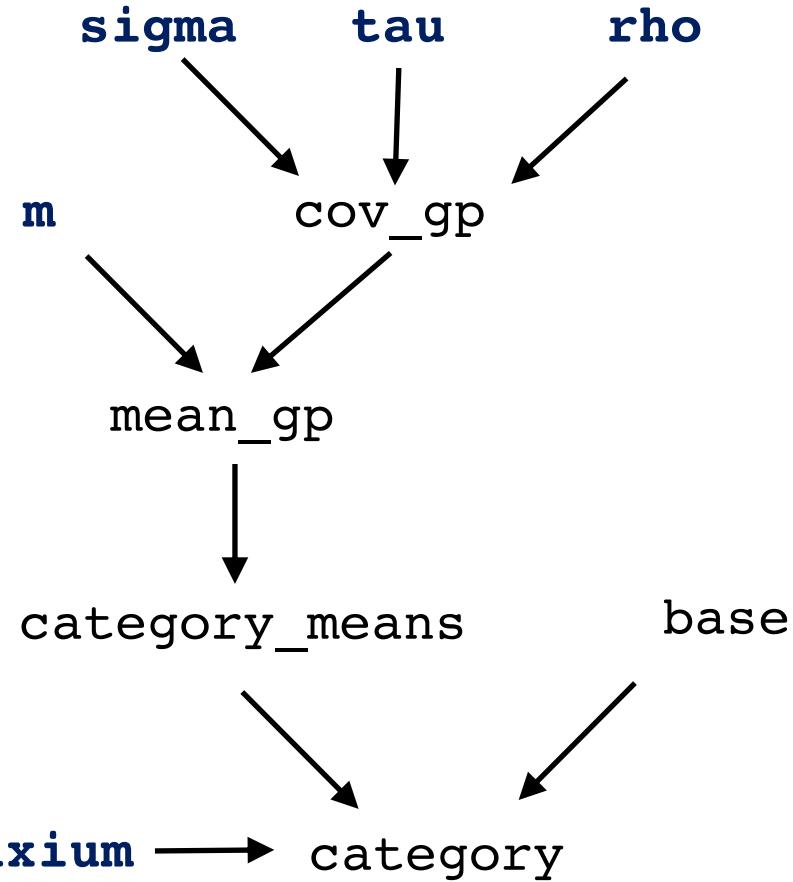
Prior on category means



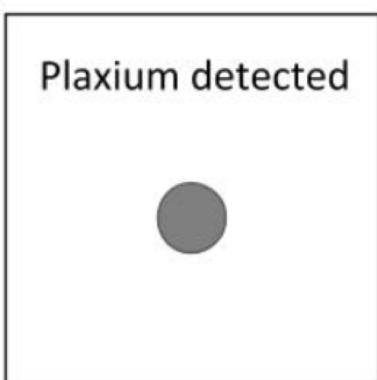
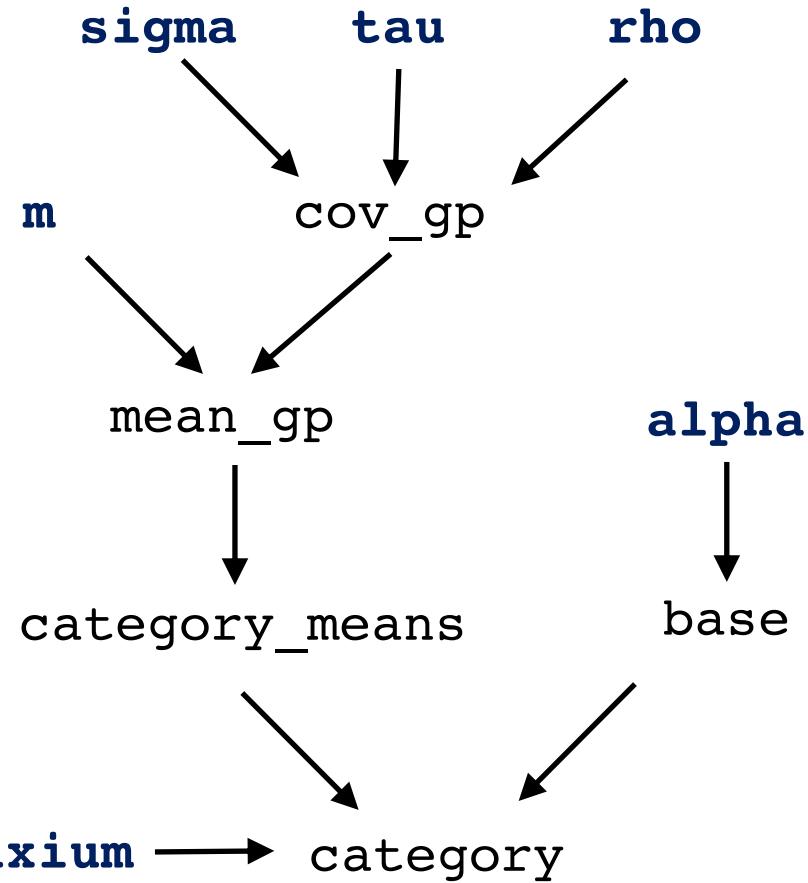
Property sampling



Property sampling

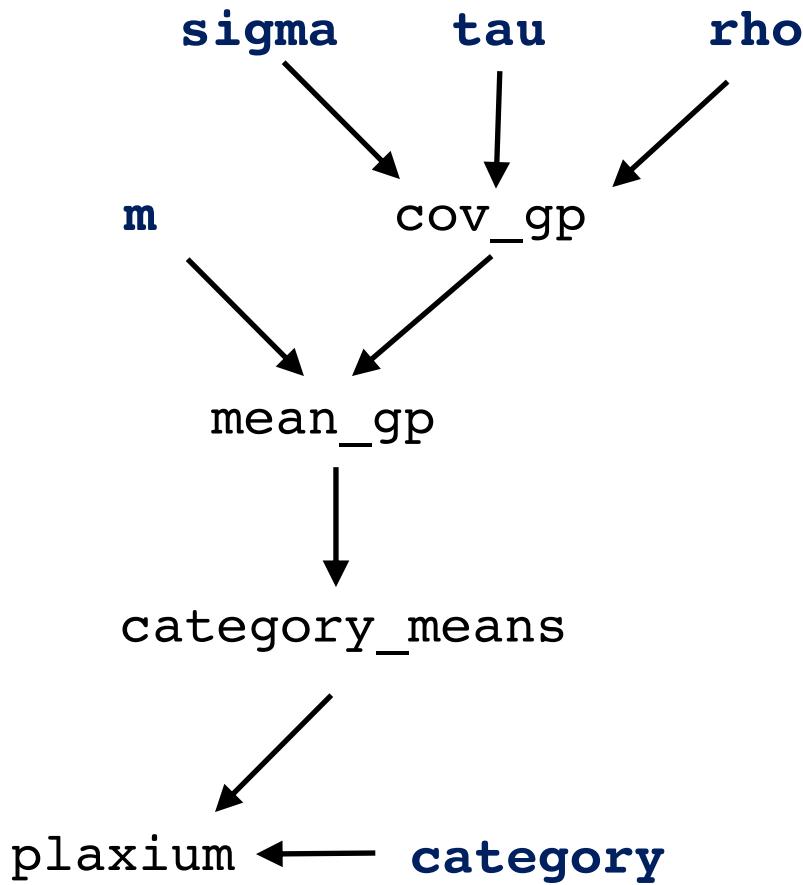


Property sampling



Inference with JAGS

Category sampling



category.bug

```
model {  
  
    # mean and covariance matrix definition  
    for(i in 1:ncat) {  
        mean_gp[i] <- m  
        cov_gp[i,i] <- (sigma^2) + (tau^2)  
        for(j in (i+1):ncat) {  
            cov_gp[i,j] <- (tau^2) * exp(-rho * (i-j))  
            cov_gp[j,i] <- cov_gp[i,j]  
        }  
    }  
  
    # sample a function from the Gaussian  
    cov_gp_inv <- inverse(cov_gp)  
    f ~ dnorm(mean_gp, cov_gp_inv)  
  
    # pass f through logistic function to get probability  
    for(i in 1:ncat) {  
        category_means[i] <- 1/(1+exp(-f[i]))  
    }  
}
```

Exercise: Sampling frames

Take home messages

- Bayes nets let you build relatively complex models out of simple pieces
- JAGS makes the implementation process relatively painless

TAKING STOCK

Today's models

- How are the models we've discussed useful?
- And what are their limitations?

Bayesian models in general

- Can Bayesian models be useful for thinking about problems you're interested in?
- Where might they be not so useful?